

Thesis Title

Politecnico di Milano



POLITECNICO

MILANO 1863

Carlos Blasco Andrés

June 26, 2020

Abstract

Handling Big Data appears as a new problem for our society that aims at extracting useful information from them.

To this purpose, complex platforms have been conceived that support the development of so-called Data Intensive Applications (DIAs). Creating a DIA, though, still requires a complex design process that, besides the problems related to ensuring high performance and availability, includes also the need to provide guarantees in terms of data privacy and to allow the users to define and update privacy policies based on their own preferences.

Thus, StreamGen is a model-driven approach to support the design of DIAs and the automated code generation for two target platforms, Flink and Spark. Moreover, StreamGen is designed with the goal to be easy to handle by users with few knowledge about DIA.

The purpose of this thesis is to extend StreamGen to make it privacy policies-aware. In order to reach such purpose: a) the definition of a language, as part of a UML profile, is required in order to allow users to define for a DIA two types of privacy policies: View Creation Policies (VCP) and Data Subject Eviction Policies (DSEP). VCPs modify the incoming data taking into account the data subject and some predefined conditions which have to be specified. DSEPs remove the data from the stream when the data comes from a data subject and it satisfies some predefined conditions. Such language is defined taking into account the dataflow model of DIAs and its sources, transformations and sinks approach, making the definition sufficiently simple to be handled by users who are not familiar with DIAs. b) from the defined language and by means of Acceleo, an implementation of a model-to-code transformation will be developed allowing users to automate the generation of the privacy-enhanced DIA code, targeting Flink as final platform. c) the evaluation of the approach will be done by exploiting two case studies. In the first case, a DIA will be developed in order to preserve the privacy of the users who make different transactions among a fixed shop stock. Secondly, a DIA will be modified. This DIA takes the temperatures from two rooms and it computes two statistics, the maximum and the average temperatures, and the prediction of the temperature of the rooms in a given time. This application will be modified in order to avoid the computation of the prediction when the temperatures come from the second room.

Resumen

El manejo de Big Data aparece como un nuevo problema para nuestra sociedad que tiene como objetivo extraer información útil de ellos.

Para este propósito, se han concebido plataformas complejas que soportan el desarrollo de las llamadas Aplicaciones Intensivas en Datos (AID). Sin embargo, crear un AID todavía requiere un proceso de diseño complejo que, además de los problemas relacionados con garantizar un alto rendimiento y disponibilidad, también incluye la necesidad de proporcionar garantías en términos de privacidad de datos y permitir a los usuarios definir y actualizar políticas de privacidad basadas en sus propias preferencias.

En este sentido, StreamGen es un enfoque basado en modelos para soportar el diseño de AID y la generación automática de código para dos plataformas de destino, Flink y Spark. Además, StreamGen está diseñado con el objetivo de que sea fácil de manejar por usuarios con pocos conocimientos sobre AID.

El propósito general de la presente propuesta de Trabajo de Final de Máster (TFM) es extender StreamGen para que tenga en cuenta las políticas de privacidad. Para alcanzar tal propósito: a) se va a requerir la definición de un idioma, como parte de un perfil UML, para permitir a los usuarios definir para un AID dos tipos de políticas de privacidad: políticas de visión del creador (View Creation Policies, VCP) y políticas de evicción del sujeto del dato (Data Subject Eviction Policies, DSEP). Las VCPs modifican los datos teniendo en cuenta quien es el propietario del dato y algunas condiciones predefinidas que se deben satisfacer sobre el mismo. Las DSEPs extraen los datos de un flujo dado cuando el dato es propiedad de un cierto sujeto y, además, satisface unas condiciones predefinidas. Dicho lenguaje se define teniendo en cuenta el modelo de flujo de datos de las AID y su enfoque en fuentes, transformaciones y sumideros, lo que hace que la definición sea lo suficientemente simple como para ser manejada por usuarios que no están familiarizados con las AID. b) A partir del lenguaje definido y por medio de Acceleo, se desarrolló una implementación de una transformación de modelo a código que permitirá a los usuarios automatizar la generación del código AID con privacidad mejorada, apuntando a Flink como plataforma final. c) Se va a realizar la evaluación del enfoque mediante la explotación de dos casos de estudio. En el primer caso, se desarrollará una AID para preservar la privacidad de los usuarios que realizan diferentes transacciones sobre un stock fijo de una tienda. En segundo lugar, se modificará una AID que toma las temperaturas de dos habitaciones y que calcula un estudio estadístico de las temperaturas máximas y medias, y la predicción de la temperatura de las habitaciones para un tiempo dado y se añadirán políticas de privacidad para que no se calcule la predicción de la temperatura de la segunda habitación.

Sommario

La gestione dei Big Data appare come un nuovo problema per la nostra società che mira a estrarre informazioni utili da essi.

A tal fine, sono state concepite piattaforme complesse che supportano lo sviluppo delle cosiddette applicazioni ad alta intensità di dati (DIA). La creazione di una DIA, tuttavia, richiede ancora un processo di progettazione complesso che, oltre ai problemi relativi a garantire prestazioni e disponibilità elevate, include anche la necessità di fornire garanzie in termini di privacy dei dati e di consentire agli utenti di definire e aggiornare le politiche sulla privacy in base a le proprie preferenze.

Pertanto, StreamGen è un approccio basato su modelli per supportare la progettazione di DIA e la generazione di codice automatizzata per due piattaforme target, Flink e Spark. Inoltre, StreamGen è progettato con l'obiettivo di essere facile da gestire da parte di utenti con poche conoscenze su DIA.

Lo scopo di questa tesi è estendere StreamGen per renderlo consapevole delle politiche sulla privacy. Per raggiungere tale scopo: a) la definizione di una lingua, come parte di un profilo UML, sarà richiesta per consentire agli utenti di definire per una DIA due tipi di politiche sulla privacy: Visualizza politiche di creazione (VCP) e politiche di sfratto degli interessati (DSEP). I VCP modificano i dati in entrata tenendo conto dell'interessato e di alcune condizioni predefinite che devono essere specificate. I DSEP rimuovono i dati dal flusso quando provengono da una persona interessata e soddisfano alcune condizioni predefinite. Tale linguaggio viene definito tenendo conto del modello di flusso di dati delle DIA e delle sue fonti, trasformazioni e approccio dei pozzi, rendendo la definizione sufficientemente semplice per essere gestita da utenti che non hanno familiarità con le DIA. b) dal linguaggio definito e tramite Acceleo, verrà sviluppata un'implementazione di una trasformazione da modello a codice che consente agli utenti di automatizzare la generazione del codice DIA ottimizzato per la privacy, indirizzando Flink come piattaforma finale. c) la valutazione dell'approccio sarà effettuata sfruttando due casi studio. Nel primo caso, verrà sviluppata una DIA al fine di preservare la privacy degli utenti che effettuano transazioni diverse tra un magazzino fisso. In secondo luogo, verrà modificata una DIA. Questa DIA prende le temperature da due stanze e calcola due statistiche, la temperatura massima e quella media, e la previsione della temperatura delle stanze in un dato tempo. Questa applicazione verrà modificata per evitare il calcolo della previsione quando le temperature provengono dalla seconda stanza.

Acknowledgements

I want to thank...

Contents

1	Introduction	17
2	State Of The Art	21
2.1	Data-Intensive Applications	21
2.1.1	Characteristics	22
2.1.2	Frameworks	23
2.1.2.1	Apache Hadoop	24
2.1.2.2	Spark	24
2.1.2.3	Apache Flink	25
2.2	Data Privacy	26
2.2.1	Privacy Policies in DIAs	27
2.2.1.1	View Creation Policies (VCPs)	29
2.2.1.2	Data Subject Eviction Policies (DSEPs)	30
2.3	Modeling and Metamodeling	30
2.3.1	Unified Modeling Language (UML)	32
2.3.1.1	UML Diagrams	33
2.3.1.2	UML Frameworks	34
2.3.1.3	Object Constraint Language (OCL)	34
2.3.2	StreamGen	35
2.3.2.1	DIA Generation Process	35
2.4	Discussion	37
3	Research Design	39
4	Solution Design	45
4.1	Privacy Policy Modeling Language	45
4.1.1	PrivacyProtectingStream	45
4.1.2	PrivacyPolicyPackage	48
4.1.2.1	PrivacyContextSource	50
4.1.2.2	PrivacyPolicySource	52
4.2	Privacy Policy Code Generation	53
4.2.1	Privacy Policy Codes Structure	54
4.2.1.1	Library Import	54
4.2.1.2	Privacy Policy Initialization	56
4.2.1.3	Streams Declaration	58
4.2.1.4	Streams Protection	59

5	Evaluation	63
5.1	Great Seller Privacy Aware DIA	64
5.1.1	Great Seller Source	64
5.1.2	Great Seller Transformations	67
5.1.3	Great Seller Sinks	68
5.1.4	Great Seller Privacy Enforcement	68
5.1.4.1	Protected Streams Definition	68
5.1.4.2	Privacy External Sources Definition	70
5.1.5	Great Seller DIA Limitations	71
5.2	Cool Analyst Privacy Aware DIA	72
5.2.1	Cool Analyst Sources	72
5.2.2	Cool Analyst Transformations	74
5.2.3	Cool Analyst Sinks	77
5.2.4	Cool Analyst Privacy Enforcement	77
5.2.4.1	Protected Streams Definition	77
5.2.4.2	Privacy External Sources Definition	79
5.3	Discussion	80
6	Conclusions	81
	Appendices	85
A	Papyrus Project Creation	87
B	Support Python Codes	91
B.1	Great Seller Stock Creation	91
B.2	Great Seller Input Tuples Creation	92
B.3	Great Seller Input Server	92
B.4	Cool Analyst Temperature Chamber 1 Model	93
B.5	Cool Analyst Temperature Chamber 2 Model	95
B.6	Cool Analyst Chamber 1 Server	96
B.7	Cool Analyst Chamber 2 Server	97
C	Run Configuration	99

List of Tables

2.1	Modeling Hierarchy	32
2.2	StreamGen Classes	36
3.1	Program Versions Used	43
3.2	Extended Language Abstract	43
4.1	PrivacyProtectingStream Properties	47
4.2	ProtectedStreamConfiguration Data Type	47
4.3	Privacy Context Source Properties Abstract	52
4.4	Privacy Policy Source Property Abstract	53
4.5	Applied Actions for Class Stereotypes in PrivacyPolicySources Package	57
5.1	Great Seller Data Types	65
5.2	Great Seller Stock	67
5.3	PrivacyProtectingStream Great Seller StreamS2	69
5.4	PrivacyProtectingStream Great Seller StreamS3	70
5.5	Cool Analyst Data Types	73
5.6	PrivacyProtectingStream Cool Analyst TempTupleStream	78
5.7	Cool Analyst External Sources Properties	79
5.8	Cool Analyst SCV Text File	80

List of Figures

2.1	Data Intensive Application Scalability	23
2.2	Modeling Method Components	32
2.3	UML Metamodel	33
3.1	Great Seller Dataflow Model	40
3.2	Great Seller Privacy Dataflow Model	41
4.1	PrivacyProtectingStream UML Profile	46
4.2	Privacy and Non-Privacy Aware Stream	46
4.3	PrivacyPolicyPackage UML Profile	48
4.4	PrivacyPolicyPackage Example	49
4.5	PrivacyPolicyDataSources Classes UML Profile	49
4.6	Dataflow Model Context Stream Generation	51
4.7	PrivacySources Example	52
5.1	Great Seller Data Types Package	66
5.2	Great Seller StreamGen DIA	69
5.3	Great Seller Privacy Aware StreamGen DIA	70
5.4	Great Seller Privacy Policy Package	71
5.5	Cool Analyst Data Types Package	74
5.6	Temperature Chamber 1 Model	75
5.7	Temperature Chamber 2 Model	75
5.8	Cool Analyst StreamGen DIA	78
5.9	Cool Analyst Privacy Aware StreamGen DIA	79
A.1	New Papyrus Project	87
A.2	Architecture Context	88
A.3	Project Name Definition	88
A.4	Representation Kind	88
A.5	Papyrus Project Definition	89
C.1	Java Project Creation	99
C.2	Java Project Initial Structure	100
C.3	Java Project Final Structure	100
C.4	Maven Project Structure Creation	100
C.5	Maven Project Source Folder Creation	100
C.6	Maven Project Properties	101
C.7	Maven Project Structure	101
C.8	POM File Creation	101
C.9	Final Maven Project	102

C.10 Run Configuration 102

Chapter 1

Introduction

The development of Information Technologies (IT) during the last century has generated new necessities in our society. In 1936, Alan Turing presents the first machine which is able to compute anything computable ([1]). Since then, researchers have been developing computing technologies until arriving to the current situation where technologies are able to handle large amount of data, terabytes and petabytes ([2]).

At the beginning, computing machines were not able to process a program with more than one instruction. Users had to introduce manually each of the instructions composing the program. This processing mode is known as interactive processing and there was not any user interface ([3]). Some time later, batch computing revolutionized computer science field. This new processing type was able to compute programs with more than one instruction automatically and, in addition to that, batch computing introduced the firsts parallel pipelines ([3]). Currently, computers are able to use stream processing which goes one step ahead than batch processing. Stream processing is able to query data continuously, computing data just when it is received ([4]).

Due to the fact that data calculation capacity has not stopped growing during the last century, currently more and more data are produced by the processors of our digital devices such as mobile phones, laptops or tablets among others. This has given rise to the well-known Big Data which is being exploited by many companies in order to extract useful information from the generated data of our digital devices.

To this purpose, complex platforms have been conceived that support the development of so-called Data Intensive Applications (DIAs). Creating a DIA, though, still requires a complex design process that, besides the problems related to ensuring high performance and availability, includes also the need to provide guarantees in terms of data privacy and to allow the users to define and update privacy policies based on their own preferences.

In order to guarantee data privacy, some laws has been legislated by public administrations in the last years. In the European Union (EU) the General Data Protection Regulation (GDPR), according to its web page [5], is in charge of harmonizing data privacy laws across all member states of the EU. However, some legislation in this regard is not enough to guarantee data privacy in digital devices. This is why, some researches, such as [6], about how privacy policies can be applied in data-intensive applications have been made during the last years.

On the other hand, the development of DIA cannot be faced writing codes from scratch. This is why complex platforms based on metamodeling techniques have

been conceived in order to reach efficiency in the DIA design process. This is the case of StreamGen, a model-driven approach to support the design of DIAs and the automated code generation for two target platforms, Flink and Spark. Moreover, StreamGen is designed with the goal to be easy to handle by users with few knowledge about DIA. However, StreamGen does not provide a way to automate code generation for privacy-aware DIAs.

Then, the purpose of this thesis is to extend StreamGen to make it privacy policies-aware. In order to reach such purpose:

1. The definition of a language, as part of a UML profile, is required in order to allow users to define for a DIA two types of privacy policies: View Creation Policies (VCP) and Data Subject Eviction Policies (DSEP). VCPs modify the incoming data taking into account the data subject and some predefined conditions which have to be specified. DSEPs remove the data from the stream when the data comes from a data subject and it satisfies some predefined conditions. Such language is defined taking into account the dataflow model of DIAs and its sources, transformations and sinks approach, making the definition sufficiently simple to be handled by users who are not familiar with DIAs.
2. From the defined language and by means of Acceleo, an implementation of a model-to-code transformation will be developed allowing users to automate the generation of the privacy-enhanced DIA code, targeting Flink as final platform.
3. The evaluation of the approach will be done by exploiting two case studies. In the first case, a DIA will be developed in order to preserve the privacy of the users who make different transactions among a fixed shop stock. Secondly, a DIA will be modified. This DIA takes the temperatures from two rooms and it computes three statistics, the maximum, the minimum and the average temperatures, and the prediction of the temperature of the rooms in a given time. This application will be modified in order to avoid the computation of the prediction when the temperatures come from the second room.

Regarding to the solutions provided along this document in order to reach the purpose of the paper, the definition of the language is based on a new data stream stereotype which allows DIA developers to define which streams are protected and which privacy policy types (VCP and DSEP) protect them. In addition to this stereotype, the language is extended with two more stereotypes which are represented by means of two generalizations of a data source stereotype. One of the extended data source stereotypes introduces the privacy rules that must be applied in order to guarantee data privacy whilst the second stereotype introduces the viewer of the DIA who is not allowed to see the streams flowing through the application without being encrypted by the privacy policies. Finally, a package where the two new data source stereotypes must be introduced is presented in order to make the approach of the external privacy sources more intuitive for the DIA developer.

Finally, it is important to remark how the document is structured and which is the content of each of the chapters. Chapter 2 consists of three main sections. Firstly, an approach to data-intensive applications is made giving the main characteristics of such applications and the most important frameworks that support

them. Secondly, the main concepts about data privacy, such as privacy policies, are introduced making a distinction on data security and data privacy. Finally, modeling and metamodeling techniques for software engineering is presented giving the most popular modeling languages and an introduction to a metamodel platform for DIA design such as StreamGen.

After that, in chapter 3, some literature of researches in privacy policy is presented and a problematic in the modeling of such researches is identified. Then, the objectives of this paper are stated and under which requirements they must be satisfied. Finally, in this chapter is presented how the developed approach is evaluated.

In chapter 4 the development of the approach is made giving an example in order to make the approach more intuitive for the reader. Moreover, the approach is classified in two parts. The part regarding to the language extension with the three main stereotypes explained above and the part referring to the Acceleo code generation in order to automatize the Flink codes generation.

In chapter 5 two case studies are developed in order to check the approach and in order to see if the objectives stated in chapter 3 are satisfied.

Finally, in chapter 6, some conclusions about the new approach are made.

Chapter 2

State Of The Art

Big Data has transformed our society in the last years. The capability to develop efficient applications which were able to handle large amount of data at run time was the first step after reaching stream processing. In this chapter the main concepts about data-intensive applications are presented by defining what they are, why data privacy has become in a must in this kind of applications and how data-intensive applications can be developed efficiently.

In section Data-Intensive Applications, first of all, a classification of these kind of application is made distinguishing between compute-intensive and data-intensive. After that, the main characteristics of data-intensive applications are explained in order to define them. Finally, in this section some frameworks as Hadoop, Spark or Flink are explained giving some characteristics and what makes them different.

In section Data Privacy some concepts relative to computer security and how such concepts impact on data privacy are developed. After that, some distinctions between data security and data privacy are made and, at the end of the section, it is emphasized how privacy policies can be applied in the already known data-intensive applications.

Finally, in section Modeling and Metamodeling, first of all the main concepts of metamodeling are exposed. After that, the Unified Model Language is explained and, at the end of the section, a model-driven approach to support data-intensive application design called StreamGen is introduced.

2.1 Data-Intensive Applications

Data-Intensive Applications (DIAs), according to [2], are a class of parallel computing applications which handle high amount of datasets and that devote the most of their processing time to I/O and manipulate such datasets. These kind of applications are able to manage datasets of several terabytes and petabytes which are available in a wide variety of formats and that are distributed among several locations. This capability to process large volumes of data is possible due to the fact that DIAs use multi-step analytical pipelines with transformations and fusion stages in order to parallelize the execution of the data.

Intensive applications can be classified ([7], [8] and [9]) according to the parallel processing approaches commonly known as compute-intensive and data-intensive. On the one hand, compute-intensive is used to describe application programs that are compute bound. Such applications devote most of their execution time to com-

putational requirements and they process small volumes of data. In this kind of applications, processes must be parallelized by means of individual algorithms and then they must be decomposed into separate tasks. Later, such tasks are executed in a pipeline reaching higher performance than in serial processing. Finally, compute-intensive applications are able to perform multiple operations simultaneously as they allow task parallelism.

On the other hand, data-intensive is used to describe application programs that are I/O bound ([10]). Such applications devote most of their execution time to I/O, move and manipulate the datasets. This second parallel processing approach, unlike compute-intensive, processes large volumes of data. Data-intensive applications require the partition or subdivision of the initial dataset in some smaller datasets due to the fact that the large volume of data that they are composed. These partitions must be processed independently by the application allowing the parallelization of the initial dataset. Once the partitions are executed, they must be reassembled in an output dataset.

As any type of application which takes advantage of data-parallelism, DIAs need to select an algorithm not only to distribute data among the processing nodes of the cluster but also to execute the data in such processing nodes ([11]). Moreover, they need to adopt a strategy to decompose the datasets, find a trade-off in terms of working loads in the processing nodes by means of load balancing systems and to communicate the nodes of the cluster.

2.1.1 Characteristics

Data-intensive applications are modeled in a different way than any other computing application due to the property of DIAs to be processed in a cluster infrastructure. Moreover, as any other system, DIAs need to reach a high performance in order to make efficient applications; this is possible by means of the minimization of the number of movements of the data. Other important properties required by DIAs are that they must be available when they are called, and that they must be reliable to the number of failures; both properties are accomplished due to data-intensive applications are fault resilient. Last but not least, data-intensive applications must be able to change the number of required nodes for processing purpose taking into account the workload. Then, four characteristics, according to [12] and [13], can be defined in order to distinguish data-intensive applications from any other kind of computing system:

Dataflow Model

Data-intensive applications are modeled by means of nodes where some transformations take place and data or programs flow between such nodes. Moreover, such model allows to control how to schedule, to execute, to balance the load and to communicate programs and data during runtime using a cluster computing infrastructure.

Data Movement Minimization

Minimize the number of movements required by each data along the different computations in order to achieve high performance is crucial in data-intensive applications.




Figure 2.1: Data Intensive Application Scalability

Then, processing algorithms are executed on the nodes of the dataflow model where the data is located, reducing the number of movements and therefore increasing performance.

Fault Resilient

Due to the fact that DIAs work with high amounts of data, the probability for the hardware to fail is greater than in other computing systems but the same happens with the probability of communications errors and software bugs. This is why, DIAs are designed to be fault resilient. In order to make these applications available and reliable, on the storage disk of such systems redundant copies of all data files or intermediate processing results can be found. But also they are provided with systems capable to detect the failure of any node or processing failures and, once such failures are detected, then re-compute the results.

Scalability

Depending on the hardware architecture of the data-intensive application, the required number of nodes for processing and the working nodes along the run time can be fix or variable. Due to the fact that the workload of a system, for example a database, can vary over time, systems can be scaled for higher processing speed. A system can be scaled horizontally or vertically. Taking into account the database example, the hard disk capacity can be raised reaching a database with more capacity. However, a database can be seen its size increased by adding more hard disks with lower capacity reaching a final database with higher capacity.

In the figure 2.1 can be seen how vertical scalability works. In this figure can be seen how a database with a big capacity increases its capability by adding a small capacity. Moreover, in the same figure can be seen how horizontal capability works as well. Some databases with small capability can be connected in parallel obtaining a unique database with big capacity.

2.1.2 Frameworks

Nowadays, several architectures are able to deal with data-intensive applications but, before arriving to the current situation, processing frameworks have been developed by means of different approaches.

At the beginning, the very first type of processing was the interactive processing ([3]). In this processing type there was not any user interface. Users had to control program executions in a computer by writing the commands manually.

After that, batch processing appeared. Batch processing, according to [3], allows to plan the work execution and execute them without user interaction. This was the first approach to parallel processing. However, batch processing requires to store data in a database, file system or any other storage unit before executing it. After data storage, data is queried for execution.

The most optimized approach for processing is stream processing. This approach allows to the queries exist continuously, computing on data just when it is received or produced ([4]).

Currently, batch processing and stream processing are the most used processing approaches all around the world. Moreover, there are a lot of frameworks which are able to handle them. However, the most important framework for such processing approaches are:

- Apache Hadoop.
- Spark.
- Apache Flink.

2.1.2.1 Apache Hadoop

Apache Hadoop, according to [14], is an open source software framework that supports distributed applications. As any other DIA, Hadoop allows to work with thousands of nodes and petabytes of data. This framework is composed by a processing layer called MapReduce which was inspired by Google documentation ([14]).

MapReduce works as a DIA processing large volumes of data. First of all, this programming model divides the incoming works into a set of independent tasks and, after that, such tasks are executed ([14]). Then, the user sends the complete job to a master which divides it into independent tasks and submits them to the slaves in order to process the works in parallel ([14]). This process allows Hadoop to reach speed and reliability of cluster ([14]).

In summary, Hadoop MapReduce works breaking the incoming datasets into independent sets that are executed in parallel by means of two phases ([14]): map phase and reduce phase. Firstly, in the map phase, the business logic is specified by putting the custom code in key-value pairs. And, secondly, in the reduce phase, aggregations are applied taking into account the key of the key-value pairs generated in the first phase.

2.1.2.2 Spark

Spark is an open source platform that allows general-purpose and fast cluster computing. It is an engine that allows to process large volumes of data.

Spark enables users to access repeatedly to datasets in order to perform streaming, machine learning or SQL workloads by means of high-level APIs in Java, Scala, Python and R. It also allows users to accomplish batch processing or stream processing.

Moreover, Spark can be integrated with any other data-intensive tool due to its design. Spark is compatible with Hadoop as it is able to access Hadoop data sources and run on Hadoop clusters. This platform extends MapReduce, the engine of Hadoop, by including iterative queries and stream processing.

Some features make Spark more efficient than Hadoop. The following features, that are presented in [15], make Spark being the 3G of Big Data:

- Fault tolerance.
- Dynamic in nature.
- Reusability.
- Advanced analytics.
- Lazy evaluation.
- Real-time stream processing.
- In-memory computing.

The in memory computing feature reduces the number of read-write to disk operations what achieves a data processing performance 100 times faster in memory and 10 times faster on the disk ([15]). Due to the high amount of operators that Spark provides, to achieve parallel applications is easier with Spark. Regarding to fault tolerance, Sparks handles the failure of the nodes in the cluster by means of Resilient Distributed Datasets (RDDs) which its specific design makes to reduce the number of failures to zero ([15]).

In conclusion, Spark can be differentiated from Hadoop in its cluster management system. On the other hand, regarding to the similarities between Spark and Hadoop, Spark takes advantage of Hadoop for storage.

2.1.2.3 Apache Flink

Apache Flink is an open source platform which allows data streaming computation because of its data flow engine. Moreover, as happens with Spark, Apache Flink is able to execute stream processing and batch processing and it is also compatible with Hadoop.

Apache Flink can undertake efficiently some different types of processing, becoming in the most potent open source platform in the market to handle them. Such processing types, according to [16], are:

- Batch Processing.
- Interactive processing.
- Real-time stream processing.
- Graph Processing.
- Iterative Processing.
- In-memory processing.

Apache Flink reduces the complexity that other platforms as Spark faced by means of some improvements. Such improvements include the integration of query optimization in MapReduce processing layer, some concepts from database systems and efficient parallel in-memory and out-of-core algorithms. These optimizations are given due to the fact that Flink architecture is designed taking into account the streaming model. Such design improves the micro-batch approach taken in Spark for data streaming processing making Flink faster and with lower latency for such kind of processing.

Following the example of Spark, Apache Flink is provided with three APIs (DataStream, DataSet and Table APIs) in order to use its engine. DataStream and DataSet APIs are two regular programs used by Flink in order to perform transformations on data streams (filtering, aggregating, update state...) and data sets (joining, grouping, mapping...) respectively. On the other hand, Table API is used for relational operations and it can be integrated into DataStream API for relational stream processing and into DataSet API for relational batch processing.

Finally, it is important to remark the five features that can be found in [16] about Apache Flink:

1. Low latency and high performance.
2. Fault tolerance.
3. Memory management.
4. Iterations.
5. Integration.

All these characteristics are what make Flink the 4G of big data and the most powerful tool when dealing with data stream processing.

2.2 Data Privacy

The appearance of Big Data and DIA platforms dealing with large amount of datasets has led the society to question about the privacy of those data that are processed. Due to the fact that such data are not visible for the data producers but these data are bought and handled by third parties in order to compute some statistics, data privacy has become in a must for the development and commercialization of new DIA in order to give the option to the user to decide on which data can be seen by other agents.

Data privacy and data security can be treated interchangeably. However, they are two terms with some differences according to [17]. First of all, data security protects data from external attackers or malicious insiders, also known as black hat hackers. On the other hand, data privacy governs how data must be collected, shared and used.

Every computer system that ensures data security, it must accomplish the well-known CIA paradigm, as it is explained in [19]. "C" stands for Confidentiality, "I" for Integrity and "A" for Availability and they are the three basic requirements that compose such paradigm. Any computer security system has to grant that any piece

of data has to be accessed only by those who are authorized, accomplishing such requirement, information is confidential. Moreover, such systems have to be able to allow data modifications only to those how are authorized for it and only in the way that they are entitled to modify them, this requirement makes information to be integrated. Finally, data must be available to everyone who has a right over them within some time constraints, this makes information to be available. These three requirements raise an engineering problem and it is that availability conflicts with confidentiality and integrity due to the fact that if a data is available then it could not be confidential as anybody may modify it and it could not be integrated as anybody could modify then.

Data privacy can be classified, according to [17], as the branch of data security concerned with the proper handling of data. Due to the fact that the information of many entities (persons, companies, etc.) are involve in such complex problem, there are some legal instruments to handle the design of data privacy. This is the case of privacy policies. According to [18], privacy policies, in the context of Information Technologies (IT), is a document that tells readers how a technology or other product or service will use their personal information. Privacy policies can differ from one country to another. This is why there are some agreements in order to handle them. In the European Union (EU) the General Data Protection Regulation (GDPR), according to its web page [5], is in charge of harmonizing data privacy laws across all member states of the EU. But GDPR is not the unique privacy regulatory restriction referring to data privacy. Some other regulations pointing to different fields can be found. Such are the cases of Health Insurance Portability and Accountability Act (HIPAA, [20]), Gramm-Leach-Bliley Act for US financial industry (GLBA, [21]) or California Consumer Privacy Act (CCPA, [22]) which is the first consumer privacy act in the US.

In addition to these regulatory restrictions, some technical tools are required in order to apply privacy policies from an engineering point of view. Some researchers haven been studying about the development of formal languages in order to specify them in a logical way. This is the case of some literature [23] where XACML is used to specify privacy policies by means of some logic rules in order to reach a set of policies. This formal language (XACML) also can be used for other reasons as it is the case of user authentication in distributed systems [24].

In a similar way that formal languages have been developed in order to specify privacy policies, some researches have been focus on applying privacy policies to dataflow applications. This is the case of [6] where two different privacy policies for DIA are developed from a metamodeling perspective.

2.2.1 Privacy Policies in DIAs

Due to the fact that DIAs are constantly managing data from different entities, it is very important the way in which privacy policies are applied on them. The final goal of applying them is that non-allowed third parties can handle those sensitive data that users do not want to be accessible by them.

When dealing with DIAs, sensitive data are sent in data streams. Such data streams are composed of tuples which can contain a high variety of data. Each tuple of a data stream contains a finite set of data that are always ordered in the same way. Due to the fact that all the tuples contained in a data stream have the

same structure, when speaking about data streams, each of the data of the tuple is considered as a field of the data stream. Due to this distinction on the fields of the data streams, some literature [6] classifies data streams into three different sets:

- Subject-specific streams.
- Subject-generic streams.
- Non-personal streams.

A subject-specific stream is a stream composed of a field referring to the data subject. This means that each tuple of the stream contains a data which specifies who is the owner of the tuple. A subject-generic stream is a stream without data subject. However, subject-generic streams are produced by an operator whose input is a subject-specific stream. Finally, a non-personal stream is a stream without data subject and that it has not been produced by an operator whose input is a subject-specific stream.

The context in which a tuple is conveyed is also important. Depending of the context in which data are transferred, an entity could want that the data of the tuple are private or public. This context is specified by means of three variables in [6]:

- Observer.
- Role.
- Purpose.

The first variable is the observer of the data, the entity who demands the usability of the data. The second variable is the role that the requesting entity is playing. The requesting entity can be a person but also a company. In the case of the company, many persons working for such company can be demanding the data of the tuple and each person has a role inside the company, this is why the role is an important variable in order to define the context in which a privacy policy can be applied. Finally, the third variable that defines the context is the purpose. The purpose for which the requesting entity demands the data of the tuple. Then, the context in which a privacy policy has to be applied is defined by means of three variables that are called static context variables (SCV) as, once they are defined, they are not going to change and because of this they are static.

Another important distinction is the one given by the data that compose a tuple. Each of these data can be strings or numbers. Because of this, a classification on the type of each data of the tuple has to be defined. Each of the data that compose a tuple are known as fields of the stream as well. Moreover, following the approach given by the SCVs, each of these values are considered as a variable since the values of a field of a given stream can change depending on the tuple that is referred. Then, the variables that compose a field of a stream, according to [6], are:

- Categorical variables.
- Numerical variables.

Categorical variables are those that correspond to a string in any programming language. They can be words but also links or any set of letters with special characters and numbers. The second type of variables correspond to numbers but also to a range which corresponds to a pair of numbers. For example, a data subject can be a categorical variable when it is the name of the user who uses the DIA but a data subject can be also a numerical variable if it is an ID which is composed only by numbers. Furthermore, the set of categorical and numerical variables used by a DIA are called contextual variables.

Taking into account these three classifications, privacy policies can be defined by means of some rules. These rules have to be defined making a relationship between the contextual variable used by the DIA and some values that the user selects in order to encrypt the original fields of the stream. These relationships are made by means of logical operators. In the case of categorical variables, the operators that can be used are: equal to ($=$) and not equal to (\neq). On the other hand, in the case of numerical variables, the relation operators that can be used are: equal to ($=$), not equal to (\neq), is greater than ($>$), is less than ($<$), is greater than or equal to (\geq) and is less than or equal to (\leq).

Finally, due to the fact that privacy policies encourage the anonymity of the entity who owns the data of the tuple, depending on the data stream classification explained above and the rules defined by each user, privacy policies can differ. Then, another classification is made in [6], in this case about the privacy policies in DIAs:

- View Creation Policies (VCPs).
- Data Subject Eviction Policies (DSEPs).

2.2.1.1 View Creation Policies (VCPs)

View Creation Policies (VCPs) are a type of privacy policies for DIAs which are applied to subject-specific streams. VCPs are applied when the context is precisely specified by means of the SCVs and the information contained in all the fields of the stream is observable, which means that it has not been modified before by another privacy policy.

Furthermore, VCPs are defined by means of a set of rules which are relationships between the contextual variables used by the DIA and some values that the user selects in order to encrypt the original fields of the stream. In this kind of privacy policies, the encrypted final values of the stream are defined by means of generalization vectors. A generalization vector defines the value to be used in order to generalize each of the fields that compose a subject-specific stream.

In conclusion, given a tuple belonging to a subject-specific stream, if at any instant such tuple satisfies all the rules defined by means of logical operators relating some contextual variables with a value defined by the user of the DIA, a generalization vector is applied on the tuple in order to generalize each of the fields of the streams and obtaining an encrypted tuple. Then in [6], a VCP is defined by means of four items:

1. Subject-specific stream.
2. Data subject.

3. Set of relational rules.
4. Generalization vector.

2.2.1.2 Data Subject Eviction Policies (DSEPs)

Data Subject Eviction Policies (DSEPs) are applied when given an operator, whose input is a subject-specific stream, produces a subject-generic stream. Sometimes, the entity who owns the data which are flowing along the DIA does not want that an operator takes some information from the input stream in spite of the output stream is not showing any information directly pointing to such owner. This is the case in which DSEPs are applied.

Unlike what happens with VCPs, DSEPs do not encrypt the fields of a given data stream. In this case, DSEPs prevent the tuples flowing in the subject-specific stream that inputs into a given operator from being processed by such operator. In this case, the output from the operator is computed without taking into account the data of the data subject who defines the DSEP, preserving the anonymity of the owner of the data for the computation of a certain statistic.

In conclusion, in order to define a DSEP, first of all, a set of rules have to be specified by the user relating some contextual variables with some values that such user has to define. Then, given an instant when all the defined rules are satisfied on a subject-generic stream which outputs from a given operator, all the tuples referring to the data subject who defined such rules have to be evicted from the subject-specific stream that inputs into the given operator. Then in [6], a DSEP is defined by means of three items:

1. Subject-generic stream.
2. Data subject.
3. Set of relational rules.

2.3 Modeling and Metamodeling

Data Intensive Applications are used when large amount of data have to be handle with high performance. In order to reach such high performance, DIAs require complex designs and a lot of resources to achieve their goal. One of these resources is the number of software developers who are needed to write the application codes. On the other hand, when some specific requirements are added to the DIAs, as it is the case of privacy policies managing the privacy of the users, developing DIAs becomes in a really difficult task which cannot be faced writing codes from scratch. At this point is where other approaches to write application codes take advantage. Such is the case of modeling and metamodeling techniques applied to software development.

A model is a representation of a reality in order to explore, to redesign or to transform it ([25]). This representation can be externalized by means of a diagram in a paper or in a computer software ([25]). Modeling is the method that has to be developed in order to reach a model. A modeling method consists of a modeling technique and a mechanism or algorithm working on the model ([26]). Moreover, a modeling technique consists of a modeling language and a modeling procedure ([26]).

A modeling procedure is a sequential process composed of several steps but also an iterative process since the procedure can be repeated many times until the final model is reached ([26]). In spite of several modeling procedures can be accomplished in order to create a model, as it can be found in [25], all these approaches share some common steps:

- Purpose definition.
- Boundaries definition.
- Model elements definition.
- Relationships definition.

The first step that has to be accomplished in order to create a model is to define its purpose. With such purpose, the questions that the model will have to answer, once it is implemented, should be specified. Later, the boundaries of the model must be specified. Such boundaries must focus on the reality that has to be modeled discarding any information that is not able to reach the predefined purpose. Once these boundaries are defined, everything that is going to take part of the model is called the model domain. Due to the fact that usually this domain is still very big, each concept belonging to the model domain must be filtered taking into account the relevance of each of them and grouping those concepts that are identical in order to reach the purpose. After that, each of the groups with identical concepts can be transformed into a model element, which is an abstraction of all the concepts that are part of the group. Finally, the relevant relationships among the model elements have to be identified and represented in the model.

Furthermore, a model must be complete in order to be able to answer the question specified in the purpose but also, it must be consistent which entails the lack of contradictions in its representation ([26]). In order to reach such characteristics, modeling methods take advantage of modeling techniques in addition to a mechanism, an algorithm, working on the model.

The first component of a modeling technique that is a modeling language is a set of well-known elements and relationships by means of which a model can be represented ([26]). A modeling language is described by a syntax, semantics and a formal notation ([26]). A syntax is the set of elements, relationships and rules that can be written in order to represent a model by means of a grammar. Moreover, a syntax can be described by two different types of grammars: graph grammars and metamodels. According to [26], a metamodel is a model of a modeling language. A metamodel is written by means of a modeling language that is known as the metamodeling language. The metamodeling language has a model which defines it. This is why exists a modeling hierarchy which is not limited and that ends given a useful abstraction level of the model. In the table 2.1 there is an example of a modeling hierarchy with five levels.

The semantics of a modeling language describe the meaning of the language by means of a semantic domain and a semantic mapping.

Finally, the formal notation describes how the modeling language is visualized. Some symbols are used in order to represent the syntax of the language in a diagram.

On the other hand, the finality of the mechanism is to check that the model written with the modeling language by means of the modeling method is correct

Language Level	Models	Language Name
Level 1	Model	Modeling Language
Level 2	Metamodel	Metamodeling Language
Level 3	$Meta^2$ -Model	$Meta^2$ -Modeling Language
Level 4	$Meta^3$ -Model	$Meta^3$ -Modeling Language
Level 5	$Meta^4$ -Model	$Meta^4$ -Modeling Language

Table 2.1: Modeling Hierarchy

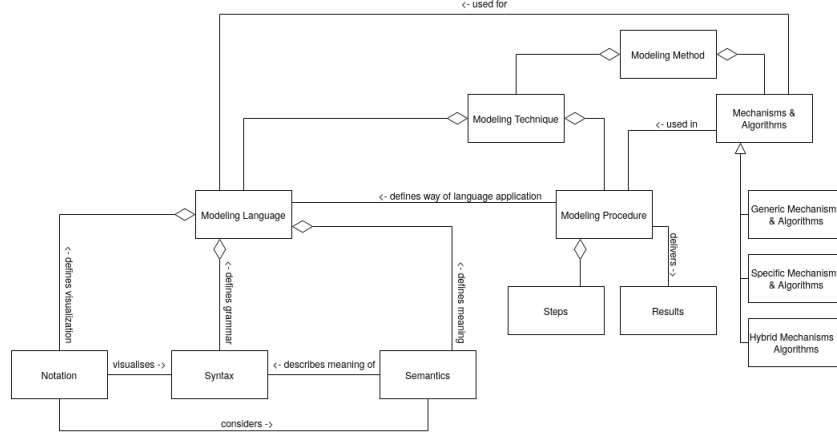


Figure 2.2: Modeling Method Components

and satisfies the purpose of the model. There are three main types of mechanisms and algorithms: generic, specific and hybrid.

In the figure 2.2 o model of a modeling method is shown. In this figure can be seen that a modeling method is composed of a modeling technique and a mechanism or algorithm. The mechanism or algorithm can be of three different types: generic, specific and hybrid. On the other hand, a modeling technique is composed of a modeling language and a modeling procedure and a classification can be made about both components. The modeling language is composed of a notation, a syntax and a semantics and the modeling procedure is composed of some steps.

2.3.1 Unified Modeling Language (UML)

Unified Modeling Language (UML) is a standardized modeling language whose objective is to provide system architects, software engineers, and software developers with tools for analysis, design, and implementation of software-based systems as well as for modeling business and similar processes ([27]).

UML is able to achieve such goal by means of a wide variety of options to build modeling diagrams. Among the fourteen possibilities that UML provides to construct a diagram, a distinction between two main classes is made in [27]: behavior diagrams and structural diagrams.

Behavior diagrams represent the dynamics of a system ([27]), this means the representation of all the possible changes that a system can experience over time. This class is composed of seven diagrams ([27]): activity diagram, communication diagram, interaction overview diagram, sequence diagram, state machine diagram,

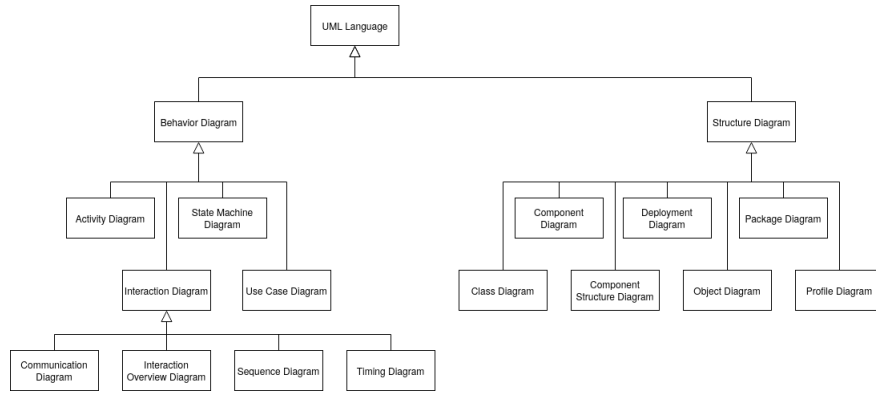


Figure 2.3: UML Metamodel

timing diagram and use case diagram.

On the other hand, structural diagrams represent the different static parts of the system ([27]), that ones which do not change over time. This representation can be done from different abstraction and implementation levels and making the corresponding relationships between the parts of the system. This second class is composed of other seven possibilities for diagram construction ([27]): class diagram, component diagram, component structure diagram, deployment diagram, object diagram, package diagram and profile diagram.

A UML diagram making a description of all the UML diagrams can be seen in the figure 2.3.

2.3.1.1 UML Diagrams

In spite of UML allows developers to use fourteen diagrams for modeling, some of such diagrams are the most commonly used. According to [27], this is the case of the class diagram and the profile diagram.

Class Diagram

This kind of diagram represents the model elements and the relationships that exist between such elements. The model elements, which are represented by means of nodes, can be: class, component, data type, interface, model or package among others. Regarding to the relationships between the nodes, they can be: association, association class, dependency, generalization or information flow among others.

Class diagrams are commonly used to represent the metamodel of the syntax of a modeling languages. Due to the fact that this kind of diagrams is not able to fully express syntactical rules, usually they take advantage of constraint languages as Object Constraint Language (OCL).

Profile Diagram

A profile diagram allows to create the model domain by adding model elements called stereotypes and making the corresponding relationships between them. The most commonly used relationships in this kind of diagrams are the composition and the generalization.

2.3.1.2 UML Frameworks

UML can be found in a wide amalgam of platforms due to its utility for modeling. However, there are some platforms that are the most commonly used for software development. This is the case of Eclipse. Eclipse is the most popular Java Integrated Development Environment which allows a large amount of plugins in order to increase its functionalities. Moreover, among such functionalities can be found UML modeling or converting a UML model to a source code.

For UML modeling one of the plugins that can be integrated in Eclipse for such purpose is Papyrus which is an open-source UML tool. On the other hand, to convert a UML model to a source code Acceleo appears as an open-source plugin for Eclipse.

Papyrus

Papyrus is one of the plugins provided by Eclipse, an open-source UML tool for UML modeling. In spite of Papyrus was developed by the French Alternative Energies and Atomic Energy Commission (CEA-List), currently it is one of the most used UML tools for software development as it can be expanded with UML profile diagrams giving a wide functionality for modeling software applications.

Acceleo

Regarding to platforms for converting an UML model into source code, Acceleo is an open-source generator from Eclipse Foundation. Acceleo is one of the Eclipse plugins that allows to transform UML models into source code files. It is written in Java and it is available for Linux, Windows and Mac OS.

Acceleo uses its own language in order to generate the source codes. Such language is based on MOF Model to Text Transformation Language (MOFM2T) and on template focusing. A template contains some text which describes the data that should be extracted from the elements of the model and such data are extracted iteratively from each component of the model. Moreover, the descriptions of the information that should be taken from the model elements are written in OCL language.

Finally, it is import to remark that Acceleo can extract the information from different types of models, such as EMF, UML or DSL and it is able to generate different source languages such as C, Java or Python.

2.3.1.3 Object Constraint Language (OCL)

Object Constraint Language (OCL) is a language to describe formal expressions in UML models. These expressions represent invariants, preconditions, postconditions, initializations, guards, derivation rules, as well as queries to objects in order to determine its state conditions.

Initially, OCL was developed by IBM. But, then, in 2003, OCL was adopted as part of UML 2.0 by the group OMG.

This language does not cause side effects, so the verification of a condition, which presupposes an instantaneous operation, never alters the objects of the model. Its main role is to complete the different artifacts of the UML notation with formally expressed requirements.

In conclusion, OCL associates boolean expressions with the elements of the model and these expressions are evaluated as true every time the program is executed. In addition, they can also be associated with methods. In this case, the Boolean expressions are evaluated as true at the moment before or after the execution of the program.

2.3.2 StreamGen

StreamGen is a model-driven approach to support the design of DIAs and the automated code generation for two target platforms, Flink and Spark. This approach is developed focusing on two requirements:

- To provide a platform-independent and graphical modeling language for streaming applications.
- To enable the fast prototyping of distributed streaming applications over different platforms.

The first requirement is needed in order to develop any kind of DIA for different target platforms and, independently for which platform it is targeted, provide a complete language for all of them. This is possible due to the DIA dataflow model characteristic that supports the development of them by means of sources, transformations and sinks. The second requirement is satisfied by using UML class diagrams for such purpose. This kind of diagrams allow to represent sources, transformations and sinks but also the data flows really fast, just introducing the predefined stereotypes.

2.3.2.1 DIA Generation Process

The different steps that have to be followed in order to generate a DIA with StreamGen are the following:

- UML model creation.
- Application code generation.
- Application execution.

UML model creation

First of all, the user must create a UML model by exploiting a predefined language in a UML profile diagram. This language consists of four metaclasses:

- Class.
- InformationFlow.
- Model.
- Package.

Sources	Transformations	Sinks
Text File	Reduce	Text File
Socket	Map	CSV File
Collection	Window	Socket
Kafka	Flatmap	Cassandra
-	NFlatmap	Kafka
-	NMap	-
-	Count	-
-	Fold	-
-	Sum	-
-	Filter	-
-	Join	-
-	CoGroup	-

Table 2.2: StreamGen Classes

The class metaclass represents the streaming operators stereotypes which are data sources, data sinks and data stream transformation. Moreover, this three stereotypes can be represented by means of several generalizations. In the table 2.2 can be shown such stereotypes.

The InformationFlow metaclass represents the data streams. There are three main data streams stereotypes: non-parallel stream, windowed stream and partitioned stream. Moreover, the partitioned stream stereotype can be generalized to broadcasted stream, keyed stream, randomly partitioned stream and round robin stream.

The model metaclass is extended by a distributed streaming application stereotype which is generalized to Flink application, Spark application and Apex application.

The last metaclass, the package, is extended by the stream data types stereotype which will contain all the data types conveyed by the different information flows.

Application code generation

Once the user has defined the UML model, such model must be transformed into some application codes. This is possible because of the development of an Acceleo code. This Acceleo program is able to take as inputs the generated UML model and, depending on the different stereotypes contained in such model, generate a DIA code for Flink or Spark.

Application execution

Once the UML model is transformed into Flink or Spark codes, the application is ready to be executed. For such purpose, first of all, the data sources are prepared and the codes are sent to the master server in order to distribute the processing tasks. This last step is the same one that any code using a batch or a streaming platform for processing.

2.4 Discussion

The appearance of Big Data has generated a growth in IT research. Firstly, with the irruption of Hadoop framework, the firsts DIAs were developed. Since then, more development frameworks, as Spark or Flink, have been improving Hadoop features until arriving to the targeted stream processing where data is computed just when it is received or produced.

Due to the fact that currently data-intensive applications provide speed besides performance, usually we can find them in our day to day and, with the generalization of DIAs in our lives, two new challenges in IT research appeared.

The first one is relative to the privacy of the large amount of data that DIAs are handling. Since DIAs are used in several fields and some of these fields, as it is the case of the bank industry, work with sensitive data, data must be modified in order to keep the privacy of the data owner what means to preserve the data privacy. Some regulations in this regard have been legislated in the last years but they are not enough. In addition to the legislation some technical approaches in order to apply the regulations have been studying. This is the case of privacy policies as View Creation Policies and Data Subject Eviction Policies.

Regarding to the second challenge, coding data-intensive applications from scratch is no longer efficient. At this point is where techniques as metamodeling appear to have an important role in the new software development scenario. The well-known UML language used by software developers to avoid programming from scratch can be used for such purpose. Moreover, there are already some approaches in this regard. This is the case of StreamGen, a model-driven approach to support the design of DIAs and the automated code generation for two target platforms, Flink and Spark. However, StreamGen does not provide a way to automate code generation for privacy-aware DIAs.

Chapter 3

Research Design

In the last years the interest among data-intensive technologies and applications in order to extract useful information from data has increased in our day to day. These data-intensive applications and their execution environments allow users not only to handle Big Data but also to do it efficiently. This is why Big Data has become in a common technology in our lives.

Furthermore, the appearance of Big Data has generated some needs in our society. One of these needs is privacy policies. DIAs require some tools which are able to preserve the privacy of users which are generating such data. These tools can vary from international laws that legislate in this regard to technical approaches that allow to encrypt the flowing data.

Due to the development of dataflow applications and the wide range of fields where Big Data is still getting involved, writing DIA codes from scratch became in an inefficient task. At this point was where modeling techniques for software development became important for DIA development. StreamGen is an example of a platform used to develop DIAs from UML class diagrams by means of Papyrus. However, StreamGen does not allow to specify privacy policies.

Some approaches can be found in the literature about privacy policies applied to dataflow applications as it is the case of [6]. This approach allows to specify two different types of privacy policies: View Creation Policies (VCPs) and Data Subject Eviction Policies (DSEPs). Each of these privacy policies are determined by means of a transformation in charge of each of the policies. Moreover, a past condition checker operator is considered in order to check not only the static context variables conditions specified by the user but also some conditions about the variables flowing through the streams of the DIA which could be given in the past. Moreover, this article introduces a source generating the viewer tuples (SCV tuples) to feed the privacy operators.

In [6] a dataflow application based on an e-commerce example is presented to explain how privacy policies could be applied to dataflow applications. Such application consists of three operators to compute some real time statistics about the transactions generated by the consumers of the e-commerce company. The statistics are sent to external companies in order to provide useful information to both parts, the clients and the external companies that buy the statistics.

In order to develop this, a source, three transformations and three sinks are connected by means of streams as can be seen in the figure 3.1. First of all, the transaction tuples are introduced into the DIA by means of a socket source. This

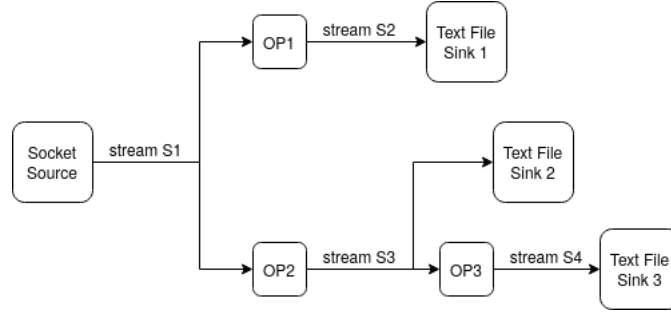


Figure 3.1: Great Seller Dataflow Model

transaction tuples flow generating the stream S1 which goes from the socket source to two transformations, OP1 and OP2. Some statistics are computed in the OP1 and the results of such statistics are sent by means of the stream S2 to the first sink, Text File Sink 1. Moreover, OP2 performs some statistics as well and, then, the results are sent to a transformation (OP3) and to a sink (Text File Sink 2) where the results flowing through the stream S3 are stored. In the OP3 transformation, some statistics are computed on the stream S3 and the results are sent to a sink (Text File Sink 3) by means of the stream S4. It is important to remark that the figure 3.1 represents the dataflow model of the application without any privacy policy.

Once the application is developed, in [6], an algorithm is presented in order to generate the privacy-aware dataflow application. Such algorithm is able to generate two privacy policy types, VCPs and DSEPs. Furthermore, the article develops two privacy policies, one VCP and one DSEP, which are applied to the e-commerce example:

- VCP: the client avoids that the employees of a market consultancy (Market-Consult employees) know if the client spent amount in the last 30 minutes is bigger than €100 and if the number of issued transactions by the client is bigger than 3, in such case, the client spent amount should be modified to the biggest allowed (€100).
- DSEP: the client does not allow that data goes to one of the operators (to operator 3) when MarketConsult observes the application.

In order to do this, some specific privacy policy operators are introduced. Such privacy policy operators are:

- StaticContextSource.
- PastConditionChecker.
- ViewBuilder.
- DataSubjectEvictor.

The proposed algorithm in [6] works with these operators as follows. First of all, the algorithm takes as input the application which is not privacy-aware. Moreover, when the algorithm is called, a StaticContextSource is directly added to the input application. Such source introduces the external companies which apply for the

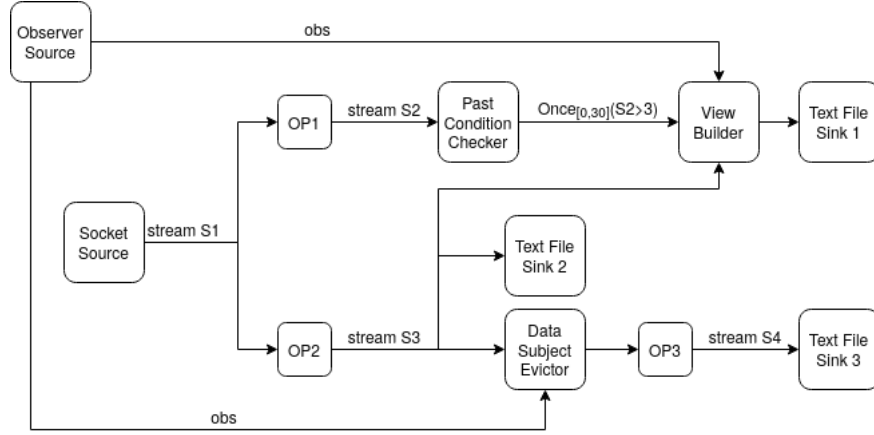


Figure 3.2: Great Seller Privacy Dataflow Model

statistical information. After that, depending on the stream, if it is VCP-aware or DSEP-aware, a ViewBuilder operator (VCP-aware stream) or a DataSubjectEvictor operator (DSEP-aware) is added to the input application. Finally, if any past condition must be taken in a given stream, a PastConditionChecker operator is applied over the corresponding stream. In the e-commerce example that is specified above, the stream S3 is VCP-aware and DSEP-aware then a ViewBuilder operator and a DataSubjectEvictor operator are applied on such stream and added to the non-privacy-aware application. Moreover, a past condition is applied for the stream S2 then a PastConditionChecker is applied on the stream S2 and added to the input application. In the figure 3.2 can be seen the privacy-aware dataflow model of the e-commerce application after applying the algorithm presented in [6].

All the Java files developed to implement the privacy-aware e-commerce application are uploaded in two GitHub repositories:

1. <https://github.com/MicheleGuerriero/dataflow-privacy-library/tree/master/src/main/java/it/deib/polimi/diaprivacy/library>
2. <https://github.com/MicheleGuerriero/common/tree/master/src/main/java/it/deib/polimi/diaprivacy/model>

Such files are stored there with the goal to provide a library to future developers to apply VCP and DSEP privacy policies in dataflow applications.

As it can be seen, this approach requires that dataflow application developers have some knowledge about the privacy policy language but also about the targeted platform language (Hadoop, Spark, Flink...) which supposes a loss of speed to design DIAs when dealing with StreamGen. These inefficiencies are given due to the fact that, after implementing the application without privacy policies, the privacy policy operators are not developed as StreamGen language and neither the translation from the UML language to the target language (Flink or Spark language). Then, there is no possibility to apply the presented algorithm in an easy way. Because of this, the current document faces the following objectives:

- To extend the StreamGen high level modeling approach to enable the modeling of privacy-aware streaming applications.

- To allow DIA developers with few knowledge about privacy policies to handle the design of privacy-aware streaming applications.
- From a non-privacy-aware application reach a privacy-aware application easily.

Thus, first of all, the StreamGen high level modeling approach, which is defined by means of a UML profile diagram, is modified in order to add all the elements which are necessary to design privacy-aware streaming applications. Moreover, the reached profile must be elegantly integrated with the current profile. In order to achieve this goal, a new language that slightly differs from the language of the four operators presented in [6] is defined. The new language identifies which are the streams that must be protected by means of a stereotype and this stereotype also has two properties to identify if the stream is VCP-protected or DSEP-protected. The extension of the language is made by means of the Eclipse plugin Papyrus. Further information about the versions used for Eclipse or Papyrus can be found in the table 3.1.

Secondly, the designer of the data-intensive application must be able to generate automatically all the Flink codes that are necessary taking into account that such designer could have few knowledge about privacy policies and Flink language. In order to reach this goal, all the Acceleo codes to compile the previously defined language are written calling to the already developed library in [6]. This allows to have the same target Flink codes independently of the DIA and, then, the Flink privacy code modification by the developer is not required. For further information about the Acceleo version used for such purpose see the table 3.1.

Last but not least, the new language must be easily applicable over non-privacy-aware applications in order to reach Flink codes without errors. This goal is achieved by means of an intuitive language based on the datatype package already implemented in StreamGen.

Finally, with the aim to see if the objectives previously stated are satisfied, the whole approach is evaluated by means of two case studies. The first case study considers the dataflow application developed in [6] for an e-commerce company. In the second case study, a DIA is developed considering an industrial application of the approach. This is the case of a factory which produces panettone. In this case study, a NMap transformation for StreamGen is developed in addition to the recognition by StreamGen of double values from the UML class diagram. The aim with both case studies is to reach Flink codes which can be run performing predefined VCPs and DSEPs.

Both cases are evaluated following the same method. First of all, the non-privacy-aware application is developed and all the statistics computed by such applications are explained. After that, the target privacy policies are explained for each case and a methodology is followed in order to apply the privacy approach and to generate the Flink codes. Finally, after applying the privacy model to the two case studies, the following questions are stated and they must be replied in order to conclude if the privacy model is complete and consistent

- Can DIA developers with few knowledge about privacy policies apply this approach?
- Can DIA developers with few knowledge about Flink language apply this approach?

- Is easily applicable the privacy approach once the non-privacy-aware DIA is already implemented?

In summary, in this document StreamGen language is extended in order to easily generate by DIA developers privacy-aware applications from non-privacy-aware dataflow application. In the table 3.2 can be found an abstract with all the added stereotypes and their metaclasses for the extension of the StreamGen language. Moreover, the Acceleo code to generate the privacy Flink codes based on the already existing privacy library presented in [6] are written. Such generated codes are defined in a way to be always the same independently of the DIA. Finally, with the aim to evaluate the developed approach, two case studies are studied, an e-commerce case study proposed in [6] and an industry 4.0 case study developed at the end of the evaluation chapter.

Program	Version
Eclipse	2019-03 (4.11.0)
Papyrus	SysML 1.4 Feature 1.3.0
Acceleo	3.7.8.201902261618
m2e-Maven Integration for Eclipse	1.11.0.20190220-2119
Apache Flink	1.4.0

Table 3.1: Program Versions Used

Stereotype	Metaclass
PrivacyProtectingStream	Information Flow
PrivacyPolicyPackage	Package
PrivacyContextSource	Class
PrivacyPolicySource	Class

Table 3.2: Extended Language Abstract

Chapter 4

Solution Design

4.1 Privacy Policy Modeling Language

StreamGen uses its own language in order to develop DIAs. In this document, StreamGen language is expanded in order to allow users to protect the streams with privacy policies. This expansion consists of a new data stream stereotype in order to specify which streams are protected and with which privacy policy type (VCP or DSEP). Moreover, the sources that generate the Static Context Variables (SCV) and the privacy rules written by the users are also added to the StreamGen language. SCVs are possibly supplied by different sources whilst privacy policies are supposed to be supplied only by one data source type. These sources are inserted in a package which represents the external sources of the DIA following the already existing approach of the StreamGen data types.

Thus, when a stream is protected, the `PrivacyProtectingStream` stereotype is added to the stream and some properties are filled in order to configure the type of protection. However, defining how the streams must be protected is not enough. Privacy policies require some external data, SCVs and privacy rules, in order to be applied. This two external data are introduced into the DIA by means of two sources, SCV source and privacy rules source. Regarding to the SCVs sources, four different types can be found in this approach. All these types are specified following the already existing approaches of data sources in StreamGen. On the other hand, for the privacy rules source, only one type is defined because the privacy rules file is supposed to be always the same. These two different sources, SCVs source and privacy rules source, are introduced in the application inside of a package stereotype called `PrivacyPolicyPackage` in order to make more intuitive the approach to the DIA developers.

4.1.1 PrivacyProtectingStream

StreamGen represents the flows of data flowing in DIAs by means of the UML metaclass Information Flow. These flows are known as streams and StreamGen allows to generate some different types as Keyed Stream or Windowed Stream. Each of these types have different characteristics and allow users to generate different behaviors. Moreover, these different streams can be applied all together generating a stream which can be, as in this case, windowed and keyed.

Following this approach, a new stereotype is added to the UML profile of Stream-

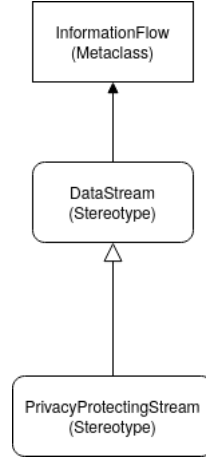


Figure 4.1: PrivacyProtectingStream UML Profile

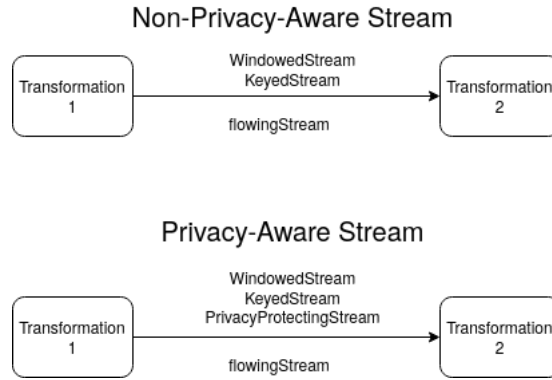


Figure 4.2: Privacy and Non-Privacy Aware Stream

Gen. This stereotype is called PrivacyProtectingStream and it is a generalization of the DataStream stereotype, which in turn is an extension of the InformationFlow metaclass. In the figure 4.1, the UML profile for the PrivacyProtectingStream is shown.

Regarding to the e-commerce company example explained in the chapter 3, let's imagine that a user defines a VCP and a DSEP in the same way that they were explained before for such example. The DIA developer, in the profile field of the property window, must specify the PrivacyProtectingStream as an applied stereotypes of the stream S3, in addition to the required stereotypes (window, keyed, etc.) for this stream.

In the figure 4.2 can be seen two streams. In the first stream the PrivacyProtectingStream stereotype is not applied in the stream called flowingStream. In this case, there is a non-privacy-aware stream. On the other hand, the flowingStream has one more applied stereotype, the PrivacyProtectingStream. In this other case, the stream is a privacy-aware stream.

Furthermore, PrivacyProtectingStream stereotype have some properties in addition to the already existing property of the DataStream stereotype, isObservable. These properties are:

- ProtectedByVCP.
- ProtectedByDSEP.

- ProtectedStreamConf.

ProtectedByVCP and protectedByDSEP properties are added in order to specify how the data stream is protected; if it is protected with a DSEP, with a VCP or if it is protected with both privacy policy types. This is why these two properties are boolean values. Moreover, the configuration of the generated protected stream, which is a stream class imported from the privacy library developed in [6], has to be specified. In order to define such configuration, a new data type has been added to the UML profile. This data type is called ProtectedStreamConfiguration and it is composed of seven properties. In the tables 4.1 and 4.2 can be seen a summary of the properties of the PrivacyProtectingStream stereotype and the properties of the ProtectedStreamConfiguration data type.

Property Name	Property Type	Multiplicity
protectedByVCP	Boolean	1
protectedByDSEP	Boolean	1
protectedStreamConf	ProtectedStreamConfiguration	1

Table 4.1: PrivacyProtectingStream Properties

Property Name	Property Type	Multiplicity
monitoringActive	Boolean	1
timestampServerIp	String	1
timeStampServerPort	Integer	1
topologyParallelism	Integer	1
simulateRealisticScenario	Boolean	1
allowedLateness	Integer	1
logDir	String	1

Table 4.2: ProtectedStreamConfiguration Data Type

Regarding to the e-commerce example again, for the stream S3 the protectedByVCP and protectedByDSEP properties would be specified as true due to the fact that a VCP and a DSEP are applied in such stream. Moreover, the values of the ProtectedStreamConfiguration data type would be specified according to the targeted protected stream that the user wants to generate.

The PrivacyProtectingStream stereotype allows to easily specify which are the streams that must be protected and with which privacy policy types. This approach does not require any knowledge about the privacy policy types. Moreover, the ProtectedStreamConfiguration data type can be filled always with the same values, then with an example using such configuration, all the other DIAs could be defined with the same values. The unique value that can vary is the logDir property which is a path and it is not difficult to be defined by a DIA developer. Another important remark that must be taken into account is that this approach does not require to specify if a past condition must be satisfied in any stream of the dataflow application. This is because of the YAML file configuration which introduces the privacy rules. Finally, this approach allows to define the non-privacy-aware dataflow application

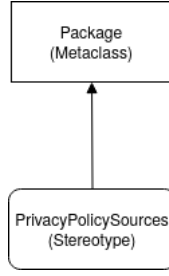


Figure 4.3: PrivacyPolicyPackage UML Profile

and, from that application, to add the corresponding privacy configurations over the streams. However, there is no way to specify the SCVs and from where the privacy rules must be read. In order to specify such sources, the `PrivacyPolicyPackage` is developed.

4.1.2 PrivacyPolicyPackage

StreamGen uses a package stereotype called `StreamDatatypes` in order to collect all the data types that are flowing through the streams. The data types that are specified inside the package allow users to specify the different values that are conveyed in a stream but inside the same tuple. Following this approach, a new package called `PrivacyPolicyPackage` is added to the UML profile in order to collect all the privacy sources involved in the privacy-aware dataflow applications, SCV source and privacy rules source.

These sources are not directly connected to the streams or to the transformations of the DIA but they allow users to specify from where the different external variables required to protect the streams are taken. As they are not directly connected to any operator of the DIA, a package to put all of them together is added to the StreamGen language in order to make the approach more intuitive for the dataflow application developers.

In the figure 4.3 can be seen the UML profile of this package. In this image can be seen how the already existing `Package` metaclass is extended with a stereotype called `PrivacyPolicyPackage` which goal is to make more intuitive for the DIA developer how to introduce the SCV source and the privacy rules source.

In order to better understand how this package would be applied in a DIA, in the figure 4.4 can be seen the `PrivacyPolicyPackage` stereotype applied in a package node and how this package node is inserted in the application model where the sources, transformations and sinks are located. In this figure can be seen that the DIA model corresponds to the the e-commerce example introduced in chapter 3.

The package where the `PrivacyPolicyPackage` stereotype is applied must be filled with two classes which correspond to the SCV tuples source and the privacy rules source. Any source corresponding to a SCV tuples source is a generalization of the `PrivacyContextSource` stereotype whilst the source corresponding to the privacy rules source is a generalization of the `PrivacyPolicySource` stereotype. Furthermore, both stereotypes are a generalization of the `DataSource` stereotype which is already developed in StreamGen. In the figure 4.5 can be seen the language of the `PrivacyContextSource` and `PrivacyPolicySource` represented in a UML profile diagram. In this figure can be seen how the already existing `DataSource` stereotype is extended

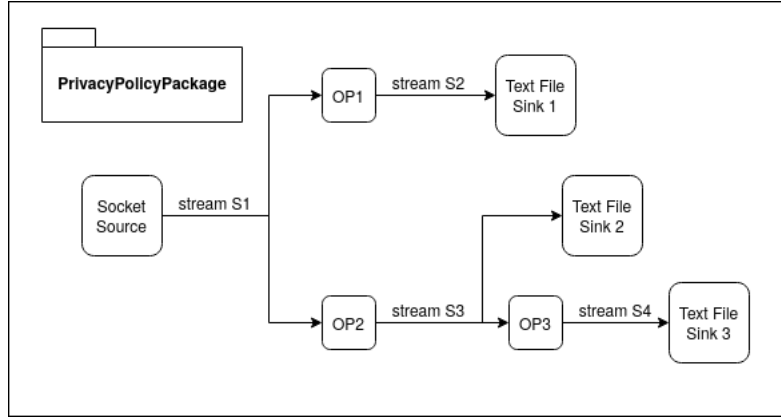


Figure 4.4: PrivacyPolicyPackage Example

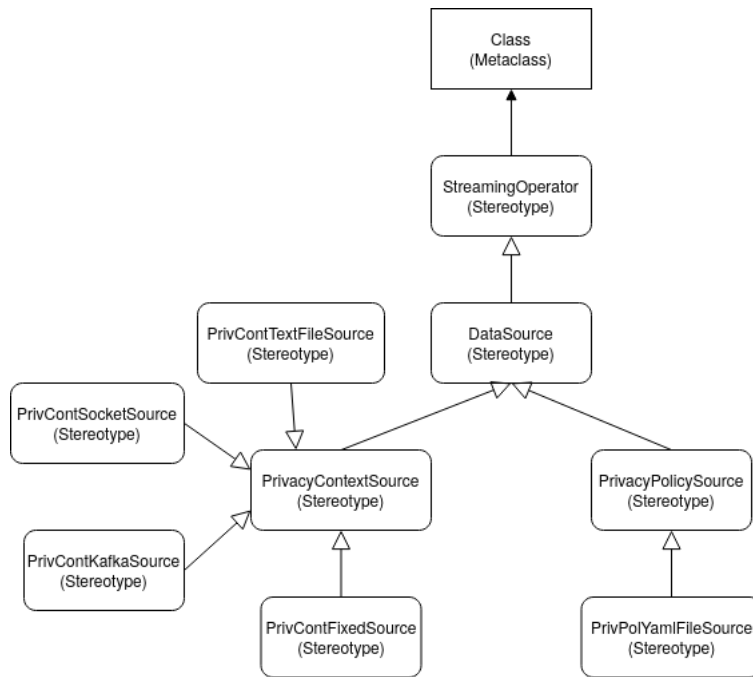


Figure 4.5: PrivacyPolicyDataSources Classes UML Profile

by means of two stereotypes, PrivacyContextSource and PrivacyPolicySource. In addition to this first extension, each of these stereotypes are extended as well. As it is explained later, the PrivacyContextSource stereotype is extended by four stereotypes (PrivContTextFileSource, PrivContSocketSource, PrivContKafkaSource, PrivContFixedSource) because the SCV source can be represented by four different type of sources. On the other hand, the PrivacyPolicySource stereotype is extended only by one stereotype (PrivPolYamlFileSource) because it can be represented only by this specific source.

In order to give an idea about the implementation of this approach, once the PrivacyProtectingStream stereotype is applied on the required streams of the DIA, a package is added to the application model and the PrivacyPolicyPackage stereotype is applied on it in order to make more intuitive to the DIA developer how the external privacy sources would be introduced later (figure 4.4).

In summary, independently of the developed DIA, two external sources are always

required. A SCV source that specifies the privacy context of the application and a privacy rules source that introduces the privacy policies defined by the users. Moreover, as these sources are not directly connected to any transformation, in order to make the approach more intuitive for DIA developers, a package called `PrivacyPolicyPackage` containing both external sources is added to the `StreamGen` language.

4.1.2.1 PrivacyContextSource

Static Context Variables (SCVs) are required in order to specify under which conditions privacy policies must be applied. According to [6], SCVs always consist on three variables: purpose, role and observer identifier. `PrivacyContextSource` stereotype provides who is observing the results of the dataflow application. Then, these SCVs can be just a tuple with the three context variables or many tuples with the three variables. This is why a wide variety of possibilities have been developed in order to allow users to choose the most appropriate SCV source depending on the DIA that is developed. Such possibilities are:

- `PrivContTextFileSource`.
- `PrivContSocketSource`.
- `PrivContKafkaSource`.
- `PrivContFixedSource`.

Following the approach of `StreamGen` with data sources, a stereotype called `PrivacyContextSource` has been created as a generalization of the already existing `DataSource` stereotype (figure 4.5). This new stereotype has four generalizations, the four possibilities among which the user can choose which is the best source for the observer depending on the application. Each of these four sources have its own properties. These properties have been defined taking into account the already existing data sources:

- `TextFileSource`.
- `SocketSource`.
- `KafkaSource`.

`TextFileSource` stereotype reads a file which contains all the input values of the DIA. When the user selects this stereotype the path (`pathToFile`, `String`) from where the DIA must read the file is specified. `PrivContTextFileSource` works exactly in the same way, this stereotype contains a property called `pathToFile` (`String`) where the user specifies the path where the file with all the SCV tuples is located. Regarding to the `SocketSource` stereotype, it is connected to a port and to a host by means of the properties `host` (`String`) and `port` (`Integer`) that the user specifies when the stereotype is defined. Following this approach, `PrivContSocketSource` is connected by a port (`Integer`) and a host (`String`) to a server which supplies the SCV tuples. Finally, the `KafkaSource` is connected by means of `kafkaBrokerIp` (`String`) and `kafkaBrokerPort` (`Integer`) to a Kafka server. `PrivContKafkaSource` connects a DIA to a Kafka server

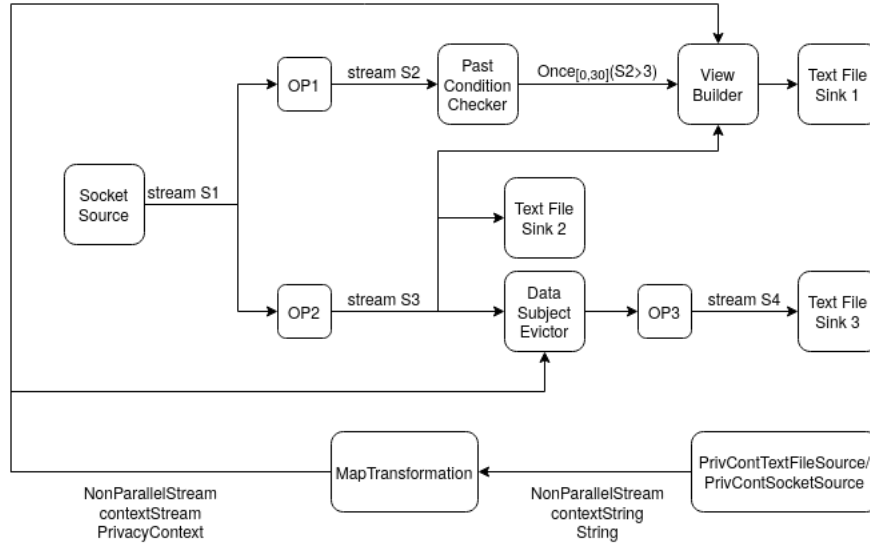


Figure 4.6: Dataflow Model Context Stream Generation

that provides the SCV tuples by means of two properties as in the case of the `KafkaSource`.

In case of `PrivContTextFileSource` and `PrivContSocketSource` stereotypes, when the code is generated, a non-parallel stream called `contextString` which contains the SCV tuples (Strings) is introduced into the DIA. This stream is automatically sent to a map transformation which is imported from the library avoiding that the DIA developer writes any Flink code. This map transformation parses `contextString` stream to a stream called `contextStream` which conveys `PrivacyContext` data types. Such data type is imported from the library in order to avoid that the application developer has to write any code and the generated stream is used to apply the privacy policies. In this document, it is assumed that independently of the developed application, these steps must be always followed. This is why when the stereotypes `PrivContTextFileSource` or `PrivContSocketSource` are applied, the Flink code relative to such stream is automatically generated. In the figure 4.6 can be seen the privacy dataflow model introduced in the figure 3.2 but the `Observer Source` is substituted by the approach proposed in this document for the stereotypes `PrivContTextFileSource` and `PrivContSocketSource`.

Finally, `PrivContFixedSource` is a source imported from the library which allows users to specify only one SCV source. In order to do this, when the user is defining the UML model of the DIA, three stereotype properties must be specified. This three properties are `fixedUser` (String), `fixedRole` (String) and `fixedPurpose` (String) which allow to specify the static or fixed SCV source.

In case of `PrivContKafkaSource` and `PrivContFixedSource` the `contextStream` stream with the user specified data type is directly generated without applying any transformation.

In the tables 4.3 can be seen an abstract with all the privacy context sources possibilities, their properties and the multiplicity of each of the properties.

In order to apply this privacy approach, after defining the `PrivacyPolicyPackage` stereotype in the corresponding package node. A class node is introduced in such package and one of the four generalization stereotypes of the `PrivacyContextSource`

Privacy Context Source	Property Name	Property Type	Property Multiplicity
PrivContTextFileSource	pathToFile	String	1
PrivContSocketSource	host	String	1
	port	Integer	1
PrivContKafkaSource	kafkaBrokerIp	String	1
	kafkaBrokerPort	Integer	1
PrivContFixedSource	fixedUser	String	1
	fixedRole	String	1
	fixesPuporse	String	1

Table 4.3: Privacy Context Source Properties Abstract

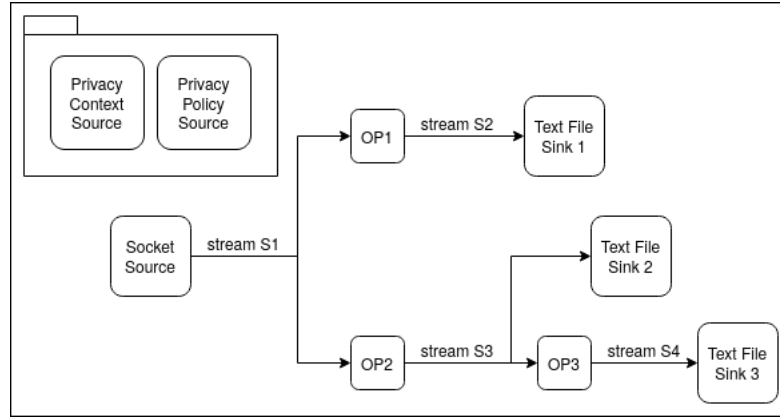


Figure 4.7: PrivacySources Example

stereotype is applied on it. Such stereotype represents the observer of the DIA, in the example of the chapter 3 this source is represented by means of a tuple, the employees of a market consultancy. In the figure 4.7 can be seen a representation for the e-commerce example where the PrivacyContextSource stereotype is applied in a class which is located inside the package with the PrivacyPolicyPackage stereotype.

4.1.2.2 PrivacyPolicySource

In addition to SCVs, StreamGen requires some privacy rules which are assumed to be specified by the users of the DIA in order to guarantee data privacy. Such privacy rules can be only some simple conditions that have to be satisfied over the stream where the privacy policy is applied or, also, some past conditions satisfied in any stream of the application. These rules are supposed to be written in a YAML file that must be inputed to the DIA in order to feed the privacy library Flink codes. In order to handle the YAML file, the library snakeyaml (org.yaml.snakeyaml.Yaml) is used. This is the unique type of source for such purpose as the rules are supposed to be specified by the user and then stored in a YAML file.

Following the approach of DataSource stereotype, a new stereotype called PrivacyPolicySource is created as a generalization of the DataSource stereotype (figure 4.5). The goal with this stereotype is to have the possibility to add different type of sources as currently is happening with the SCV source in spite of the approach developed in this document only considers one of them. Moreover, this stereotype

makes the language more understandable. From the PrivacyPolicySource stereotype, a new generalization is created in order to make the relationship with a new stereotype called PrivPolYamlFileSource. This stereotype represents the insertion of the YAML file with all the privacy policies specified by the users whose data is handled by the DIA. In order to specify the path where the YAML file is located, a new property (pathToFile, String) is added to the stereotype. With this property, the user is able to define the location where the file is stored and from where the DIA must read the privacy policy rules written by the user. In the table 4.4 can be seen an abstract of such stereotype. Following how this approach would work, after applying the corresponding PrivacyContextSource stereotype in its class, a new class would be introduced in the package with the applied PrivacyPolicyPackage stereotype and the PrivPolYamlFileSource would be applied on it. This source represents the set of privacy rules that the users define and that must be applied on the corresponding protected streams. In the example of the chapter 3 the user of the DIA would be the client of the e-commerce company. In the figure 4.7 can be seen a representation for the e-commerce example where the PrivacyPolicySource stereotype is applied in a class which is located inside the package with the PrivacyPolicyPackage stereotype.

Stereotype Name	Property Name	Property Type	Property Multiplicity
PrivPolYamlFileSource	pathToFile	String	1

Table 4.4: Privacy Policy Source Property Abstract

This stereotype follows the approach of the TextFileSource stereotype as happened with the PrivContTextFileSource stereotype. However, it is supposed to introduce a YAML file instead of a text file to the dataflow application.

4.2 Privacy Policy Code Generation

StreamGen exploits Acceleo in order to compile StreamGen language into Spark or Flink languages, depending on the target platform of the dataflow application. In this document, Acceleo is used in order to compile the StreamGen language explained above only into Flink code. Moreover, the developed privacy policy language must be translated in a general way allowing that, independently of the DIA, Flink code must not be modified by the application developer.

The above UML metamodel that specifies the language for the privacy policies must be used by the DIA developer in order to create the UML class diagram of the DIA in Papyrus. After that, this diagram is used by Acceleo in order to generate the final code which runs in Flink. In this section all the developed Acceleo codes are explained. Furthermore, all along this section an assumption is taken. This assumption is that the UML file generated by Papyrus is going to contain all the packaged classes in the following order:

1. Sources.
2. Transformations.
3. Sinks.

First of all, all the sources introducing data are packaged as elements in the UML file that is generated by Papyrus. After them, all the transformations are packaged. Finally, all the sinks where the data must be stored are packaged. This assumption is similar to say that the non-privacy-aware DIA is built in the above order and without adding transformations, sinks or sources after generating the codes of the non-privacy-aware DIA. This is important because due to the approach used for the compilation of the StreamGen language, if this order is not followed, some errors appear when the code is generated. This assumption is taken in the non-privacy-aware StreamGen version too.

In addition to this assumption, the target code generated by means of StreamGen follows a structure that generates the same pieces of code from the UML language specified before.

4.2.1 Privacy Policy Codes Structure

The generated codes for the designed privacy-aware DIA approach follow a predefined structure in order to use correctly the privacy policy libraries developed in [6]. Each of these parts correspond to a different piece of code that is generated taken into account different conditions. Moreover, as specified before, these pieces of code allow to generate the same Flink codes only with slightly modification taken into account the streams of the dataflow application allowing that the approach can be used independently of the designed DIA. This structure is defined by four parts:

1. Library Import.
2. Privacy Policy Initialization.
3. Streams Declaration.
4. Streams Protection.

In the Library Import part all the required libraries for the correct code generation operation are called in the corresponding file. Secondly, in the Privacy Policy Initialization part, all the Flink codes for the privacy sources (SCV source and privacy rules source) are generated. After that, in the third part Streams Declaration, all the streams flowing through the dataflow application are declared in order to completely fill the YAML file which is assumed to be defined by the users of the application. Finally, the streams that must be protected are protected with the privacy library developed in [6].

4.2.1.1 Library Import

First of all, a condition is required in order to import the libraries of [6] only when they are required. Then, libraries are imported when exists at least one PrivacyProtectingStream stereotype in the application model. If the DIA is non-privacy-aware, there is no PrivacyProtectingStream stereotype and the different libraries are not imported. Such libraries are imported in the file generateFlinkApplication.mtl which is where they are required for the correct code generation and the generated piece of code is the following:

```
[ if (aModel.eAllContents ( DirectedRelationship )
    ->exists (stream |
        hasStereotype (stream , ' PrivacyProtectingStream ' ))) ]
import org.apache.commons.io.FileUtils;
import java.io.File;
import org.yaml.snakeyaml.Yaml;
import it.deib.polimi.diaprivacy.library.GeneralizationFunction;
import it.deib.polimi.diaprivacy.library.ProtectedStream;
import it.deib.polimi.diaprivacy.model.ApplicationDataStream;
import it.deib.polimi.diaprivacy.model.ApplicationPrivacy;
import it.deib.polimi.diaprivacy.model.DSEP;
import it.deib.polimi.diaprivacy.model.PrivacyContext;
import it.deib.polimi.diaprivacy.library.PrivacyContextParser;
import it.deib.polimi.diaprivacy.model.VCP;
import it.deib.polimi.diaprivacy.library.PrivacyContextFixedSource;
[ / if ]
```

As it can be seen, the imported libraries are:

- import org.apache.commons.io.FileUtils;
- import java.io.File;
- import org.yaml.snakeyaml.Yaml;
- import it.deib.polimi.diaprivacy.library.GeneralizationFunction;
- import it.deib.polimi.diaprivacy.library.ProtectedStream;
- import it.deib.polimi.diaprivacy.model.ApplicationDataStream;
- import it.deib.polimi.diaprivacy.model.ApplicationPrivacy;
- import it.deib.polimi.diaprivacy.model.DSEP;
- import it.deib.polimi.diaprivacy.model.PrivacyContext;
- import it.deib.polimi.diaprivacy.library.PrivacyContextParser;
- import it.deib.polimi.diaprivacy.model.VCP;
- import it.deib.polimi.diaprivacy.library.PrivacyContextFixedSource;

The library `it.deib.polimi.diaprivacy.library` can be found in the repository <https://github.com/MicheleGuerriero/dataflow-privacy-library/tree/master/src/main/java/it/deib/polimi/diaprivacy/library>.

Regarding to the library `it.deib.polimi.diaprivacy.model`, it can be found in the repository <https://github.com/MicheleGuerriero/common/tree/master/src/main/java/it/deib/polimi/diaprivacy/model>.

Finally, the `org.yaml.snakeyaml.Yaml` library must be imported from the POM file. In order to do this, in the POM file generation `Acceleo` file (`generateFlinkPom.mtl`), a dependency is triggered when a `PrivacyProtectingStream` exists in the application model. The generated piece of code for such purpose is the following:

```
[ if (aModel.eAllContents(DirectedRelationship)
  ->exists(stream |
    hasStereotype(stream, 'PrivacyProtectingStream')))]
  <dependency>
    <groupId>org.yaml</groupId>
    <artifactId>snakeyaml</artifactId>
    <version>1.17</version>
  </dependency>
[/ if]
```

4.2.1.2 Privacy Policy Initialization

At this part, StreamGen extracts the information supplied by the user in the PrivacyPolicyPackage stereotype. StreamGen generates the corresponding SCV source Flink code and the privacy rules source Flink code. In order to do this, each of the packages contained in the application model are taken and if the package contains the PrivacyPolicyPackage stereotype, for each of the classes contained in the package a different action is applied taking into account the applied stereotype of each class. In the following Acceleo code can be seen the written generation algorithm for the privacy sources.

```
[comment StreamGen Privacy Policies Sources/]
[ if (aModel.eAllContents(DirectedRelationship)
  ->exists(stream |
    hasStereotype(stream, 'PrivacyProtectingStream')))]
  // privacy init
[ for (p:Package | aModel.eContents(Package)) ]
[ if hasStereotype(p, 'PrivacyPolicyPackage')]
  [comment]
  [generateFlinkPrivacyPolicyYamlFile(p)]
  [/comment]
  [ for (c:Class | p.eContents(Class)) ]
    [ if hasStereotype(c, 'PrivPolYamlFileSource')]
    [generateFlinkPrivacyPolicyYamlFileSource(c)]
    [ elseif hasStereotype(c, 'PrivContFixedSource')]
    [generateFlinkPrivacyContextFixedSource(c)]
    [ elseif hasStereotype(c, 'PrivContKafkaSource')]
    [generateFlinkPrivacyContextKafkaSource(c)]
    [ elseif hasStereotype(c, 'PrivContTextFileSource')]
    [generateFlinkPrivacyContextTextFileSource(c)]
    [ elseif hasStereotype(c, 'PrivContSocketSource')]
    [generateFlinkPrivacyContextSocketSource(c)]
    [/ if]
  [/ for]
[/ if]
[/ for]

// finish privacy init
[/ if]
```


The stereotypes that can be applied in the classes which are contained in the package with the PrivacyPolicyPackage stereotype are:

- PrivPolYamlFileSource.
- PrivContFixedSource.
- PrivContKafkaSource.
- PrivContTextFileSource.
- PrivContSocketSource.

In the table 4.5 can be seen an abstract with the actions performed depending on the class stereotype that is included in the package where the PrivacyPolicyPackage stereotype is applied. In such table, each of the rows represents: for each class stereotype that can be contained in the package with the PrivacyPolicyPackage stereotype, the action that is applied.

Class Stereotype	Applied Action
PrivPolYamlFileSource	generateFlinkPrivacyPolicyYamlFileSource
PrivContFixedSource	generateFlinkPrivacyContextFixedSource
PrivContKafkaSource	generateFlinkPrivacyContextKafkaSource
PrivContTextFileSource	generateFlinkPrivacyContextTextFileSource
PrivContSocketSource	generateFlinkPrivacyContextSocketSource

Table 4.5: Applied Actions for Class Stereotypes in PrivacyPolicySources Package

All the actions are applied when there is any of the class stereotypes specified in the table 4.5. Moreover, such actions are written in the generateFlinkSources.mtl file of the StreamGen Acceleo project as public templates.

In the following commented codes can be seen the examples of such generation actions for the class stereotypes PrivPolYamlFileSource and PrivContTextFileSource.

```
[comment PrivPolYamlFileSource/]
[template public generateFlinkPrivacyPolicyYamlFileSource
    (aClass : Class)]

Yaml yaml = new Yaml();

String content = FileUtils.readFileToString(
    new File("[getStereotypeProperty(aClass,
        'PrivPolYamlFileSource', 'pathToFile')/]" ),
    "UTF-8");

ApplicationPrivacy app=yaml.loadAs(
    content , ApplicationPrivacy.class );

[/template]
```

```
[comment PrivContTextFileSource /]
[template public generateFlinkPrivacyContextTextFileSource
                                                    ( aClass : Class )]

// begin privacy source definition
DataStream<String>contextString=
    env.readTextFile("[getStereotypeProperty(aClass ,
        'PrivContTextFileSource ', 'pathToFile ')]")
    [ if getParallelism(aClass ,
        'PrivContTextFileSource ').toString()
        .toInteger()>1]
        .setParallelism([getParallelism(aClass ,
            'PrivContTextFileSource ')]);
    [ else ]
    ;
[/ if ]

DataStream<PrivacyContext>contextStream=
    contextString.map(new PrivacyContextParser());
// finish privacy source definition
[/template]
```

As it can be seen in the previous code, when the source for the privacy context reading from a text file is generated, once the stream composed of strings with the three values of the privacy context is introduced in the dataflow application, it is automatically sent to a transformation in order to parse the strings into the PrivacyContext data types imported from the privacy library by means of a map transformation which generates the stream called contextStream. Moreover, the YAML file is parsed automatically into an ApplicationPrivacy class which is imported from the privacy library as well.

4.2.1.3 Streams Declaration

This document works with the YAML file configuration which is used in the repositories of [6]. This YAML file owns a field called concreteStream that is null when the YAML file is defined by the users of the DIA. This field must be filled for all the streams flowing through the application because later the field is used by the library to protect the streams in the Streams Protection part. Then, after generating and parsing the SCVs and the privacy policy YAML file in the Privacy Policy Initialization part, the generated Flink code full fills the parsed YAML file defined by the users of the dataflow application by filling the concreteStream field of each stream that is included in the file with the corresponding stream identifier.

Regardless of what stream is considered, they are non-protected once in the application. If the application does not protect a stream, the stream exists only without privacy protection. However, if the stream is protected, first of all the non-protected version of the stream is generated and, then, the stream is protected. This is why, every generated stream by the dataflow application can be declared in the YAML file. In order to do this, a condition must be taken in account every time that a stream is generated by means of a transformation. Such condition is the same condition considered in order to import the libraries; when in the application model

exists at least one PrivacyProtectingStream stereotype, then the output stream of the transformation must be declared in the YAML file. In the following Acceleo code can be seen the piece of code written for such purpose. This piece of code is replicated in all the existing transformations of StreamGen (generateFlinkTransformations.mtl).

```
[ if ( aClass.getModel().eAllContents(DirectedRelationship)
      ->exists(stream |
        hasStereotype(stream, 'PrivacyProtectingStream')) ) ]
app.getOutputStreamByID("[getOutputNames(aClass)->first()/]")
    .setConcreteStream([getOutputNames(aClass)->first()/]);
[/ if ]
```

Thus, any time that a transformation is generated, after generating the piece of code corresponding to the transformation, one condition is taken into account. This condition checks that exists at least one PrivacyProtectingStream stereotype in the application model. If such condition is satisfied, then, the stream is set into the privacy policy YAML file.

4.2.1.4 Streams Protection

StreamGen consists of three different types of operators: sources, transformations and sinks. Sources introduce some strings that must be parsed by a transformation in order to generate a data type. This happens independently of the implemented data source. Transformations make some computation modifying and generating new data types. Finally, some of the generated data types are stored by means of a sink.

Regarding to the inputs and outputs of such operators, two distinctions must be done. On the one hand, sources and sinks only have inputs (sinks) or outputs (sources) and they can have several inputs or outputs depending on the operator. On the other hand, a transformation can input a pair of streams (CoFlatMap and CoMap transformation) and, also, it can generate several streams. StreamGen is implemented in a way that assumes that all the generated streams are the same stream. Then, in spite of several streams must be specified in the UML diagram, Acceleo takes only one of such outputs in order to generate the transformation. Following this approach, privacy policies are going to be generated only once assuming that all the privacy output streams of a transformation are the same stream.

Then, the Acceleo algorithm developed in StreamGen is extended. Before extending the algorithm, StreamGen checks one by one all the transformations that are contained in the application model and depending on the applied stereotype to the classes, generates different actions. If this is not modified, it allows to generate all the streams, regardless of whether they are protected. Once the non-protected stream generated by the transformation is created, an extension of the algorithm is made. A condition is taken into account in order to check if the output stream must be protected or if it must not. In order to do this, if exists an output stream from the transformation that contains the stereotype PrivacyProtectingStream, an action called generateFlinkPrivPolicies is executed. An example only for the SocketSource, MapTransformation and TextFileSink stereotypes of the extended algorithm is presented in the following code. As it can be seen there, the protection over the streams is only applied on one of the output stream objects of the transformations. This

ensures that any stream is protected only once as it takes only one of all the replicas of the output streams. Moreover, two assumptions are taken. The first one is that any source requires a parser to transform the input stream string into a DIA data type. The second one is that all the output streams generated by a transformation are protected with the same privacy policy types (VCP, DSEP or both of them).

```
for (c:Class | aModel.eContents(Class)) ]
  [ if hasStereotype(c, 'SocketSource')]
    [generateFlinkSocketSource(c)/]
  [ elseif (hasStereotype(c, 'MapTransformation'))]
    [generateFlinkMapTransformation(c)/]
    [ if(c.getOutputs()->
      exists(output | hasStereotype(output,
        'PrivacyProtectingStream')))]
      [generateFlinkPrivPolicies(c.getOutputs()->
        select(output | hasStereotype(output,
          'PrivacyProtectingStream'))->
          first())/]
    [/ if ]
  [ elseif (hasStereotype(c, 'TextFileSink'))]
    [generateFlinkTextFileSink(c)/]
  [/ if ]
[/ for ]
```

This extended algorithm triggers the generateFlinkPrivPolicies action. Such action consists of four steps and it is written in the generateFlinkTransformations.mtl file of the StreamGen project. Moreover, this action takes advantage of the data types of the privacy library described in [6].

First of all, the stream declared in the Streams Declaration part is called and a ApplicationDataStream class imported from the library is created. After that, a protected stream with the values specified in the ProtectedStreamConfiguration data type of the property protectedStreamConf of the PrivacyProtectingStream stereotype is generated. Such protected streams is a ProtectedStream class and it is imported from the privacy library. Then, taking into account if the stream is VCP or DSEP protected with the properties: protectedByVCP and protectedByDSEP, the VCPs and DSEPs of the YAML file are added to the protected stream previously declared. Finally, all the privacy policies are applied on the stream by means of the finalize method of the protected stream. In the following code can be seen the written generateFlinkPrivPolicies action:

```
[template public generateFlinkPrivPolicies(
  stream:DirectedRelationship)

  // begin stream privacy conf
  ApplicationDataStream app_[stream.eGet('name')/] =
    app.getStreamByID("[stream.eGet('name')/]");

  ProtectedStream <[stream.eGet('conveyed').name/]>
    [stream.eGet('name')/]_p =
    new ProtectedStream
```

```

<[stream.eGet('conveyed').name/]>
([getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('monitoringActive')
    .toString()
    .toBoolean()/],
"[getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('timestampServerIp')
    .toString()/]"
[getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('timestampServerPort')
    .toString()
    .toInteger()/],
[getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('topologyParallelism')
    .toString()
    .toInteger()/],
[getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('simulateRealisticScenario')
    .toString()
    .toBoolean()/],
[getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('allowedLateness')
    .toString()
    .toInteger()/],
"[getStereotypeProperty(stream,
    'PrivacyProtectingStream',
    'protectedStreamConf')
    .eGet('logDir')
    .toString()/]"
);

[stream.eGet('name')/]_p.setStreamToProtect(
    (DataStream<[stream.eGet('conveyed').name/]>)
    app-[stream.eGet('name')/]
        .getConcreteStream());
[if isVCPprotected(stream)]

```

```
        for (VCP vcp : app.getVCPs(
            app_[stream.eGet('name')/]
                .getId())) {
            [stream.eGet('name')/]_p
            .addVCP(app_[stream
                .eGet('name')/], (VCP) vcp, app);
        }
    [/ if ]
    [ if (isDSEPprotected(stream)) ]

        for (DSEP dsep : app.getDSEPs(
            app_[stream.eGet('name')/]
                .getId())) {
            [stream.eGet('name')/]_p
            .addDSEP(app_[stream
                .eGet('name')/], (DSEP) dsep, app);
        }
    [/ if ]

    DataStream<[stream.eGet('conveyed').name/]>
        [stream.eGet('name')/]_f =
        [stream.eGet('name')/]_p
            .finalize(env, contextStream);
    // finish stream privacy conf
[/template]
```

Chapter 5

Evaluation

The evaluation of the developed privacy approach is made by means of two case studies. The aim of such case studies is to check if the stated objectives of this document (chapter 3) are satisfied. Specially, this chapter evaluates if software developers with few knowledge about privacy policies are able to handle the design of privacy-aware dataflow applications and if it is easy to reach a privacy-aware DIA from a non-privacy-aware streaming application.

In order to do that, first of all, the case studies are presented independently in two sections. Each section explains the context where the application is applied and which kind of statistics are computed. After that, the non-privacy-aware application is implemented explaining which kind of stereotypes are applied in the different sources, transformations and sinks as well as the required data types for each stream. Finally, the privacy approach is evaluated. In order to evaluate it, it is applied on the already developed non-privacy-aware application following two predefined steps:

1. Protected Streams Definition.
2. Privacy External Sources Definition.

Finally, the UML class diagram developed with Papyrus is converted into the Flink codes by running a configuration in Eclipse which executes the developed Acceleo codes. The obtained Flink codes must be generated without errors independently of the implemented application.

In summary, this section faces the evaluation of the developed privacy approach and answers the following questions:

- Can software developers with few knowledge about privacy policies apply this approach?
- Can software developers with few knowledge about Flink language apply this approach?
- Is easily applicable the privacy approach once the non-privacy-aware DIA is already implemented?

5.1 Great Seller Privacy Aware DIA

In this section a DIA for an e-commerce company called Great Seller is implemented. This example implements the case study proposed in [6]. This proposed e-commerce DIA is able to compute real time statistics about the transactions generated by the consumers. Such statistics can be observed by some other companies that pay in order to have access to the information making that Great Seller acts as a data broker. In order to simplify the implementation, Great Seller is going to sell three different types of statistics:

- Statistic 1: the total amount of money that each consumer of Great Seller spend in the last 10 minutes.
- Statistic 2: the number of transactions issued by each user in the last 10 minutes.
- Statistic 3: the number of users who spent more than €1000 in the last hour.

As any DIA, Great Seller generates such statistics by some transformations which are fed from a source and the generated information is stored into a sink which is accessible by the observer companies. Moreover, as Great Seller is producing three different types of statistics, its DIA requires three different sinks where the information should be stored. Due to the fact that Great Seller DIA is computing real time statistics about the transactions that are generated by the consumers of the company, only one source is required for the DIA model. Finally, one transformation is necessary for the computation of each statistic. In summary, Great Seller DIA requires one source, three transformations and three sinks for the design of its model.

The first step, after generating the Papyrus project (appendix A), is to define the model in the class diagram in order to specify that the application is going to be a Flink application. Then, first of all, we have to create a model node in the greatSellerApp.di file. This node is going to be called greatSeller as it is the object representing the whole application. Once the node is inserted, the properties of such node must be specified. More in detail, the Flink application stereotype has to be applied. In this case, the properties of the model has to be the default ones then nothing else must be done with this node. In the figure 5.1 can be seen how the stereotype has been applied to the model node.

5.1.1 Great Seller Source

In the implementation described along this section, Great Seller DIA is fed from a socket that reads from a text file where all the transactions generated by the consumers of Great Seller are stored. All the transactions are going to have the same predefined tuple structure 'transactionId,dataSubject,spentAmount,purchasedProduct'. This predefined structure is represented in the Great Seller DIA model by means of a data type called InputTransaction.

Regarding to the fields of such tuple, the transactionId field is an integer which varies from 1 to the number of generated transactions, such number can be 10, 100 or 1000. The dataSubject is the user who generates such tuple and, in this document, they possible data subjects are defined by a set of four names: Bob, Carlos, Elisabetta or Michele. The spentAmount is the price paid for the product

that has been bought with the transaction and it is an integer with a low boundary of €1 and an upper boundary of €200. Both boundaries are defined in this document and slightly vary from the boundaries proposed in [6] because they are considered too large for the scope of the privacy approach checker. Finally, the recipientId is the product bought with such transaction.

The steps that must be followed in order to add the data types that are required to the application are the followings. First of all, a package node has to be inserted inside the Flink application model node. This package is called great-SellerDataTypes and a stereotype has to be applied as in the previous case. First of all, inside the properties window, in the profile field, the applied stereotype is defined. Inside this package all the datatypes that the application requires have to be inserted. The first data type is the InputTransactions which is composed of a transaction id (integer), a data subject (string), an amount (double) and a recipient id (string). The second data type is the IssuedTransactions which contains the data subject (string) and the number of transactions performed by such data subject which can be seen by means of the variable NTransactions (integer). The third data type is the SpentAmount which is composed of the data subject (string) and the total amount of money spent by the data subject which can be seen by means of the variable TotalAmount (double). Finally, the last data type is the NumberUsers which contains the number of users who spent more than 1000 dollars and it is represented by means of the variable NTopUsers (integer). Then, we need to insert one DataType node inside the package created in the previous step with each of these data types. Each data type can be seen as a tuple which is composed by several values. The name of the DataType node is going to be the same that the one corresponding to the tuple and each of the values that compose the tuple are going to be a property inside of the owned attributes that are specified in the UML field of the properties of each DataType node. In the table 5.1 can be seen an abstract with all the data types required for this example. Moreover, in the figure 5.1 can be seen the package with all the data types of the application inserted in the application model with all the applied properties explained above.

Data Type	Property Name	Property Type
InputTransactions	transactionId	integer
	dataSubject	string
	amount	integer
	recipientId	string
IssuedTransactions	dataSubject	string
	nTransactions	integer
SpentAmount	dataSubject	string
	totalAmount	integer
NumberUsers	NTopUsers	integer

Table 5.1: Great Seller Data Types

It is important to remark that socket source introduces in the DIA a data stream of strings and the InputTransaction data type must be generated by means of a transformation which parses the string into the data type. This transformation is called TupleParser as it splits the incoming tuples, strings, into the InputTransaction

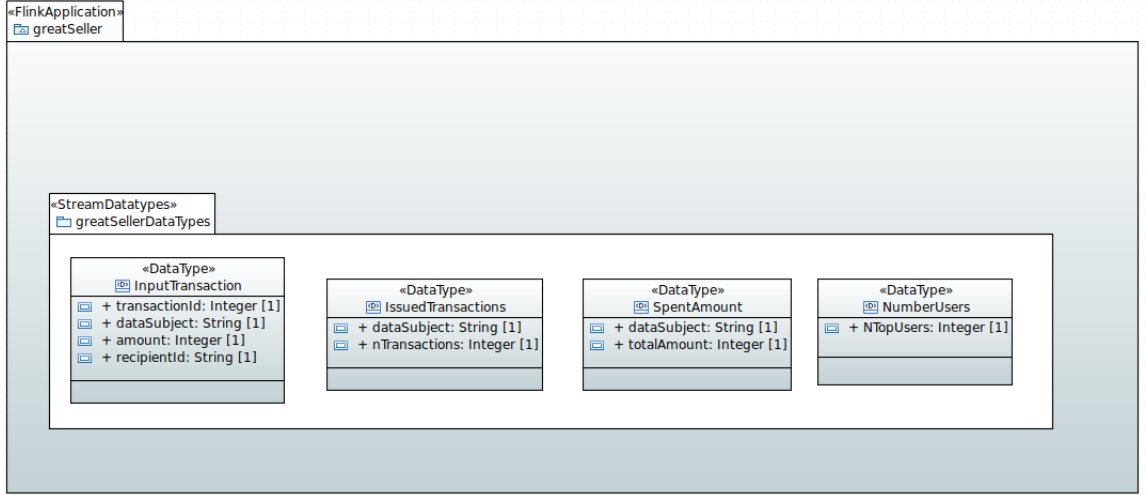


Figure 5.1: Great Seller Data Types Package

data type.

The Great Seller stock is composed by means of 25 products. In order to simplify the implementation, each product is named by the word 'product' immediately followed by a number between 1 and 25 in order to specify the product referred in the stock. In the table 5.2 can be seen such stock.

In order to generate this stock that feeds the Great Seller DIA, two Python codes have been developed. Furthermore, a Python server has been developed to input the transaction tuples to the dataflow application.

The first Python code (appendix B.1) builds a Python list which represents the Great Seller stock of 25 products and it assigns to each product a price. Each product is represented with a dictionary variable that contains a name variable and a price variable. The name variable is a string with the word 'product' immediately followed by an integer number from 1 to 25 which points to the product of the stock. The price is an integer number between \$1 and \$200 which is assigned randomly. Once the name and the price are assigned to the dictionary, the product is added to the stock list. Finally, the Great Seller stock list is saved in a binary shelf file in order to be accessible from the other Python code.

The second Python code (appendix B.2) generates the strings that represent the tuples produced by each consumer. In order to generate such tuples an integer number for the transactionId is assigned following an increasing numerical order from 1 to the maximum number of generated transactions which is input by command line to the code, it can be 10, 100 or 1000. After that, randomly, one of the four possible data subjects (Bob, Carlos, Elisabetta and Michele) and a product from the stock saved in the binary shelf file are assigned. Finally, each value is added to a string where each of these values are separated by a comma. The generated string represents the tuple produced by each consumer and it is written in the text file which is called from the Great Seller DIA in order to feed it.

Finally, the DIA inputs each of these strings into the DIA by means of a Python server (appendix B.3). This Python server introduces a non-parallel stream with the generated tuples (strings) and the stream is sent to the first transformation of the dataflow application.

Product	Price (€)
product1	196
product2	36
product3	179
product4	17
product5	120
product6	187
product7	139
product8	52
product9	160
product10	110
product11	113
product12	67
product13	100
product14	125
product15	192
product16	115
product17	113
product18	98
product19	113
product20	185
product21	143
product22	18
product23	194
product24	41
product25	26

Table 5.2: Great Seller Stock

5.1.2 Great Seller Transformations

The first transformation implemented in the application is the TupleParser. This transformation is a Map Transformation that is implemented in order to split the input strings supplied by the Python servers into the InputTransaction data type. This transformation takes advantage of the split Java method and it splits the strings by the comma generating the InputTransaction data type. The generated stream is sent to the OP1 and OP2 operators.

In order to compute each of the three statistics that Great Seller sells as data broker, the Great Seller DIA needs three transformations. Each of this transformations is the operator of each statistic. Thus, the operator one (OP1) computes the total amount of money spent by each user in the last 10 minutes. The second operator (OP2) computes the number of transactions issued by each user in the last 10 minutes. Finally, the operator three (OP3) computes the number of users who spent more than €1000 in the last hour.

These transformations input a data stream with a set of tuples that all of them have identical structure. Thus, the first stream (S1) is composed of tuples of the kind InputTransaction. This stream is duplicated and it is sent to the operators OP1

and OP2. Moreover, as each operator works taking into account the data subject of each tuple, the stream S1 must be keyed by the data subject field of this first stream. Finally, as the operators have to compute the statistics taking into account only the tuples generated in the last 10 minutes, the stream S1 is windowed by a time window of 10 minutes.

The new data generated in the OP1 are represented in a new data type called SpentAmount and that is composed by two fields following the structure: 'dataSubject, totalAmount'. The SpentAmount tuples are collected in the stream S2 and they are keyed by the data subject field of such stream.

Finally, the data generated in the OP2, is collected in a new data type called IssuedTransaction whose structure is 'dataSubject, nTransactions'. This new stream (stream S3) is sent to the OP3. Moreover, this third stream is keyed by the data subject field of the stream and it is also windowed with a time window of one hour. The OP3 generates a randomly partitioned stream with an uniform distribution called stream S4 whose structure is: 'nTopUsers'.

5.1.3 Great Seller Sinks

Each of the tuples flowing through the streams S2, S3 and S4 are stored in a file. Such values are stored because the observer companies are able to access to the tuples. This is why, after each stream is generated (streams S2, S3 and S4), a file text sink is implemented. These file are going to be stored in the same location but with different names. The file resultsS2.txt stores the generated tuples in the stream S2; the file resultsS3.txt stores the values flowing through the stream S3 and the file resultsS4.txt stores the values of the stream S4.

This implementation is given due to the fact that each of the tuples generated by the operators are observable for the observer companies who buy them. Furthermore, such sinks allow to check if the transformations are well implemented by the designer of the DIA. In the figure 5.2 the developed model with StreamGen for the non-privacy-aware Great Seller DIA can be seen. In this figure is represented the same dataflow model that in the chapter 3 but it is represented in Papyrus and it can be seen the applied stereotypes for the different sources, transformations and sinks in addition to the stream stereotypes.

5.1.4 Great Seller Privacy Enforcement

At this point, the non-privacy-aware dataflow application for the e-commerce company Great Seller is completed and the approach for the privacy-aware DIAs that is developed in this document must be applied in order to see if the previously stated objectives are met. The enforcement of the privacy model is made by means of two steps:

1. Protected Streams Definition
2. Privacy External Sources Definition

5.1.4.1 Protected Streams Definition

Once the non-privacy-aware application is implemented, the privacy metamodel is applied. First of all, the streams that should be protected are specified. In

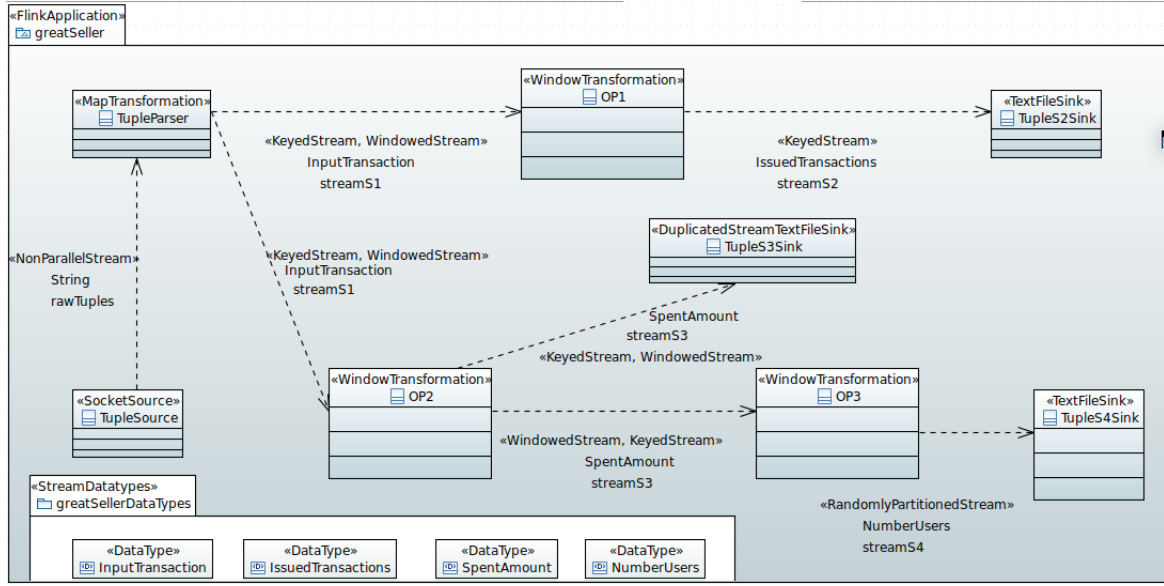


Figure 5.2: Great Seller StreamGen DIA

the Great Seller application, such streams are streamS2 and streamS3. Then, the PrivacyProtectingStream stereotype is applied on both streams. However, streamS2 is protected by a VCP whilst streamS3 is protected by a DSEP. This means that the properties of the stereotypes are different.

In the case of the streamS2 stereotype, as this stream is protected by a VCP, only such property must be true. In addition to this, the protectedStreamConf is specified according to the values that can be shown in the table 5.3. The timestampServerIp field of the protectedStreamConf is empty, this is why nothing is specified.

Property	Field Name	Field Value
protectedByVCP	-	true
protectedByDSEP	-	false
protectedStreamConf	monitoringActive	false
	timestampServerIp	
	timeStampServerPort	-1
	topologyParallelism	1
	simulateRealisticScenario	false
	allowedLateness	0
	logDir	/home/cablan/Desktop/thesis/conf/

Table 5.3: PrivacyProtectingStream Great Seller StreamS2

StreamS3 is specified in a similar way, in this case the protectedStreamConf is exactly the same but the protectedByVCP property is false and the protectedByDSEP is true. In the table 5.4 can be seen the specified. As in the other case, the timestampServerIp field is empty because no value is specified there.

In the figure 5.3 can be seen how looks the application after defining the streams that are protected. This figure represents the same model that the one that is shown in 5.2 but adding the PrivacyProtectingStream stereotype on those streams that are protected.

Property	Field Name	Field Value
protectedByVCP	-	false
protectedByDSEP	-	true
protectedStreamConf	monitoringActive	false
	timestampServerIp	
	timeStampServerPort	-1
	topologyParallelism	1
	simulateRealisticScenario	false
	allowedLateness	0
	logDir	/home/cablan/Desktop/thesis/conf/

Table 5.4: PrivacyProtectingStream Great Seller StreamS3

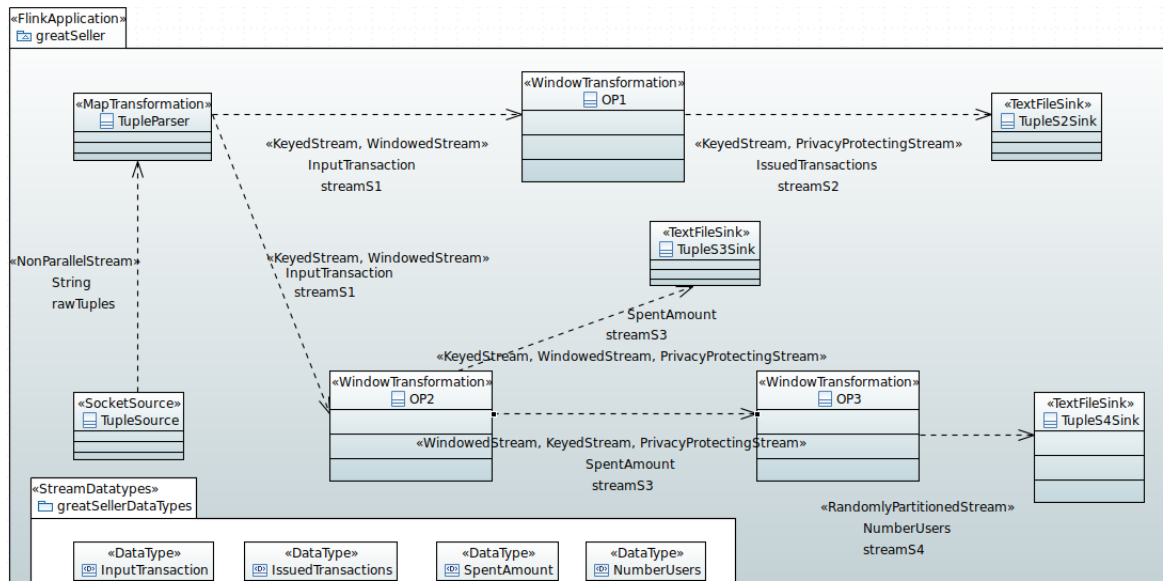


Figure 5.3: Great Seller Privacy Aware StreamGen DIA

5.1.4.2 Privacy External Sources Definition

Once the privacy streams have been defined, the second step is to define the SCV source and the privacy rules source. The privacy rules source have only one choice; however, the SCV source can be defined in several ways. In this example, as the input is working with sockets, the socket source is defined for this purpose.

But before introducing the classes of the sources, the PrivacyPolicyPackage stereotype must be defined in a new package in order to make more intuitive the approach to the DIA developer. In order to do this, a new package node is introduced in the application model node and the stereotype PrivacyPolicyPackage is applied in such new package that is called PrivacyPolicyInputs. Once the package is introduced into the model, the SCV source is defined. In order to do this, a new class is introduced inside the package called PrivacyPolicyInputs and the PrivContSocketSource stereotype is applied on it. Such new class is called StaticContextVariablesSource and their properties are filled with the values localhost (host, string) and 9998 (port, integer). Finally, another new class is introduced in the package. In this case, it is called PrivacyPolicySource and the




Figure 5.4: Great Seller Privacy Policy Package

PrivPolYamlFileSource stereotype is applied on it. This stereotype contains only one property (pathToFile) that must be filled. In this case, it is filled with the value /home/cablan/Desktop/thesisFiles/config/privacy-config.yml that is the path where the YAML file is located.

In the figure 5.4 can be seen the introduced package with the PrivacyPolicyPackage stereotype applied on it and two nested classes that represent the privacy sources with their applied stereotypes, PrivContSocketSource and PrivPolYamlFileSource. As it can be seen there, the package makes intuitive the way in which the privacy external sources are introduced to the dataflow application and the sources classes follow the same approach that any StreamGen source.

At this point, the privacy-aware DIA is implemented and the configuration is run in order to generate the Java codes of the application. In the appendix C can be seen how to run such configuration.

5.1.5 Great Seller DIA Limitations

After developing this DIA, some limitations are identified relative to the current development of StreamGen. The first limitation is the number of servers that can be connected directly to a Map Transformation. In this example, it is supposed that the users send the transactions by means of a server to the target DIA. This server is the same for all the users who are using Great Seller DIA. At this point, a casuistry is proposed. This is the case of a DIA involved in an industry 4.0 application where the owner of the factory cannot afford to invest in a server for the DIA and all the machines generating datasets must be connected directly to the DIA by means of the factory net.

The second limitation that is extracted is that StreamGen does not allow to generate float or double data types. In this example, Great Seller is working with a stock where all the products have integer prices but should be necessary a DIA which allows, among the language, to declare float or double values.

In the following example, both casuistry are fixed and the privacy model is applied to check its behavior.

5.2 Cool Analyst Privacy Aware DIA

In addition to Great Seller, another DIA is developed with StreamGen in order to check more in detail the developed approach for the privacy-aware DIAs. This is the case of Cool Analyst, a DIA which is able to compute some statistics from the temperatures measured in two different refrigeration chambers of a panettone factory called Panettone 4.0. These measurements are going to be stored in a CSV file in order to be accessible by the manager of the industrial plant to see if any problem has been produced in the yeast fermentation. This application is going to compute four statistics in two different operators:

- Operator 1: the maximum, the minimum and the average temperature of each chamber during the last 5 minutes.
- Operator 2: the prediction of the temperature for the next 10 minutes.
- Operator 3: the filtered data between a range which have no null value and no null id.

Cool Analyst DIA is going to require two socket sources, one for each of the industrial chambers, two transformations in order to compute each of the operators specified before and two sinks where the CSV files with each of the results of the operators can be stored. For the development of the prediction operator, as its development is out of this document, it is taken from an already working transformation for such purpose.

Following the Great Seller approach, before generating the dataflow application, the first step is to create the Papyrus project according to the appendix A. After that, the application model is introduced in the Papyrus project. In order to do this, a model node is defined in Papyrus and the stereotype FlinkApplication is applied on it. As in the previous example, the default properties for such stereotype are the final ones, then, they must not be modified. In the figure 5.5 can be seen the Papyrus model and the applied stereotype for the Cool Analyst application model.

5.2.1 Cool Analyst Sources

Panettone 4.0 is using temperature digital sensors in each of its two chambers where yeasts are fermented. These sensors are connected to the net of the factory and are sending the measured values directly to the Cool Analyst DIA with a 1 second sample frequency.

This two chambers need to control the temperature of the room in order to allow the correct fermentation of the yeasts. The role of the first chamber is to make the primary fermentation in the production process of the Panettone 4.0 factory. This fermentation is also known as bulk fermentation. On the other hand, the role of the second camera is to make the secondary fermentation which is also known as proofing. More information can be read about the fermentation of the yeasts and how to control the temperature during its production in [28]. The conclusions that are extracted from [28], and that are useful for our purpose, are that the bulk fermentation must have the temperatures of the room in the range [20, 24] Celsius whilst the secondary fermentation must have an environment with temperatures in

the range [22, 29] Celsius. This is why Cool Analyst requires two socket sources connected to the port and to the host of each of the chambers.

The two sockets send a tuple with the chamber identifier (room1 or room2) and the temperature of the chamber in a string separating both values by a comma. This string, as happens in the Great Seller example, is sent to a parser which generates a data type and, then, Cool Analyst works with several data types that must be defined in a package according to the StreamGen behavior. In order to do this, a package node is introduced into the model node and it is called coolAnalyst-DataTypes. After that, the StreamDatatypes stereotype is applied on the package and all the data types required by the DIA are defined. The first data type is the roomTemp data type. It is composed of two different values: roomId (string) that is the identifier of the chamber and roomTemp (double) which is the current temperature of the chamber. The second required data type is the roomStatistics data type which is composed of four values: roomId (string), maxTemp (double), minTemp (double) and avgTemp (double). The three last values of the roomStatistics data type (maxTemp, minTemp and avgTemp) are the maximum, minimum and average computed values of the chamber identified with the roomId string. Finally, the last required data type for the Cool Analyst DIA is the tempPred which is a prediction of the chambers for the following ten seconds. This data type is composed of the values roomId (string) and predTemp (double) which are the predicted temperature of the room identified by the roomId string. For each of the three data types required by the Cool Analyst DIA, a Data Type node is inserted in the package called coolAnalystDataTypes and a property for each of the values that compose the data type are inserted as an owned attribute of the data type. In the table 5.5 can be seen an abstract with all the data types required by the dataflow application. Furthermore, in the figure 5.5 can be seen the package with all the data types of the Cool Analyst application, explained above, inserted in the application model according to the StreamGen approach.

Data Type	Property Name	Property Type
roomTemp	roomId	string
	roomTemp	double
roomStatistics	roomId	string
	maxTemp	double
	minTemp	double
	avgTemp	double
tempPred	roomId	string
	predTemp	double

Table 5.5: Cool Analyst Data Types

Also it is important to remark how the temperatures of the chambers are simulated. For such purpose a python code has been developed for each of the chambers (appendices B.4 and B.5). These models take as target temperature, the average temperature of each of the ranges. This temperature is supposed to be maintained during ten seconds. Then, randomly, the program generates an increasing or decreasing sine function with a frequency of 1/16 Hz for the lower range and with a frequency of 1/28 Hz for the upper range. The lower range is considered from the




Figure 5.5: Cool Analyst Data Types Package

lower boundary of each of the ranges of temperature that the chambers can admit ([22, 29] Celsius in the chamber one and [20, 24] Celsius in the chamber two) until the average temperature and the upper range from the average temperature until the upper boundary of the ranges. In the figures 5.6 and 5.7 can be seen both models. The simulation is built for one hundred temperature values.

Finally, these two models are saved in a TXT file which is going to be read from two Python servers (appendices sec:Appendix2Sec6 and sec:Appendix2Sec7) that are developed in order to simulate the chambers working independently. These servers read one value each second during the period that the simulation lasts. This generates that the socket sending the values from the chamber one stops earlier than the other socket. This model design allows to proof the right behavior of the whole application, allowing to see what would happen if, suddenly, the sensors are disconnected.

5.2.2 Cool Analyst Transformations

The values read from the TXT file are sent in a string which has to be parsed in order to generate the data types that are used in the DIA. The strings contain the roomId followed by the temperature and both values are separated by means of a comma. This is why, first of all a NMap transformation called `generateFlinkCoMapTransformation` is developed. In order to do this, in the `generateFlinkTransformations.mtl` file, the existing but not developed `generateFlinkCoMapTransformation` is written taking as approach the transformations called `generateFlinkCoFlatmapTransformation` and `generateFlinkMapTransformation`. In the following code can be seen the written Acceleo code.

```

[template public generateFlinkCoMapTransformation(
    aClass : Class)
{
    first : DirectedRelationship =
        getInputs(aClass)->at(1);
    second : DirectedRelationship =
        getInputs(aClass)->at(2);}]

// begin stream definition

```




Figure 5.6: Temperature Chamber 1 Model




Figure 5.7: Temperature Chamber 2 Model

```

DataStream<[getOutputsConveyed(aClass)->first()/]>
    [getOutputNames(aClass)->first()/] =
[ first.eGet('name')/]
[ if first.hasStereotype('KeyedStream') ]
.keyBy("[getStereotypeProperty(first,
                                'KeyedStream', 'key')
.eGet('name')/]" )
[ elseif first.hasStereotype('BroadcastedStream') ]
.broadcast()
[/ if ]
.connect(
[ second.eGet('name')/]
[ if second.hasStereotype('KeyedStream') ]
.keyBy("[getStereotypeProperty(second,
                                'KeyedStream', 'key')
.eGet('name')/]" )
[ elseif second.hasStereotype('BroadcastedStream') ]
.broadcast()
[/ if ]
)
.map(new [aClass.name/]())
[ if getParallelism(aClass, 'NMapTransformation')
                                .toString().toInteger() > 1 ]
.setParallelism([getParallelism(aClass,
                                'NMapTransformation')/]);
[ else ]
;
[/ if ]

[ if(aClass.getModel()
    .eAllContents(DirectedRelationship)
    ->exists(stream|hasStereotype(stream,
                                'PrivacyProtectingStream')) ]
app.getStreamByID("[getOutputNames(aClass)->
                                first()/]" )
    .setConcreteStream([getOutputNames(aClass)->
                                first()/]);
[/ if ]
// finish stream definition
[/template]

```

This NMap transformation inputs the two streams with the strings sent by the room servers and the incoming strings are parsed into a roomTemp data type which is composed of two properties: roomId (string) and roomTemp(double). Furthermore, in order to introduce double values from the Papyrus model by means of StreamGen, the generateDataTypes.mtl file is slightly modified. The modification is made considering the code previously developed for the Long values. Then, every time that a property name is EDouble, it is substitute by Double. In the figure following code can be seen one of the added pieces of code to the file and how this

approach follows the Long approach. Moreover, this piece of code is added several times, as many times as necessary.

```
[for (p: Property | aClass.attribute) separator('\n')]
    private [p.type.name.substitute('ELong', 'Long')
              .substitute('EDouble', 'Double')/]
    [p.name /];
[/for]
```

Once the roomTemp data type is generated, a stream is sent to a window transformation called RoomStatistics. In this transformation all the computations explained previously in the operator one are calculated (maximum temperature, minimum temperature and average temperature). This transformation generates a new data type called roomStatistics. Then, the stream with the roomStatistics data types is sent to the TemperaturePredictor transformation. This window transformation computes the operations described in the operator two and generates the tempPred data type. The last operator takes the stream with the roomTemp data type from the RoomStatistics transformation and, by means of a filter transformation called CleanRawData, checks that there is no temperature with a value below -9999 neither 9999. Moreover, this transformation checks that the temperature has a value and that value is referenced to a room identifier.

5.2.3 Cool Analyst Sinks

Finally, the data types generated in the RoomStatistics, in the TemperaturePredictor and in the CleanRawData transformations are stored each one in its sink. Cool Analyst uses CSV sinks as this kind of files are commonly used by engineers who work in the factories. In order to do this, the output stream from the RoomStatistics transformation is sent to StatisticsSink where is specified the path of the generated file, the output stream from the TemperaturePredictor transformation is sent to PredictorSink where the path of the second output file is specified and the same procedure is done with the third operator CleanRawData and its sink CleanDataSink.

In the figure 5.8 can be seen the Cool Analyst StreamGen DIA explained above which is the non-privacy-aware application of the Cool Analyst DIA.

5.2.4 Cool Analyst Privacy Enforcement

At this points, the privacy approach for dataflow applications is applied to the already developed non-privacy-aware DIA model by defining the two already known steps. This example allows to check what happens with the developed approach for privacy-aware dataflow applications when a transformation generates a stream that is required to be protected partially, depending on the destination transformation.

5.2.4.1 Protected Streams Definition

In this DIA a DSEP is required in order to avoid tuples coming from the first chamber to be considered in the computation of the RoomStatistics transformation. From the designer perspective, no more information is required in order to define

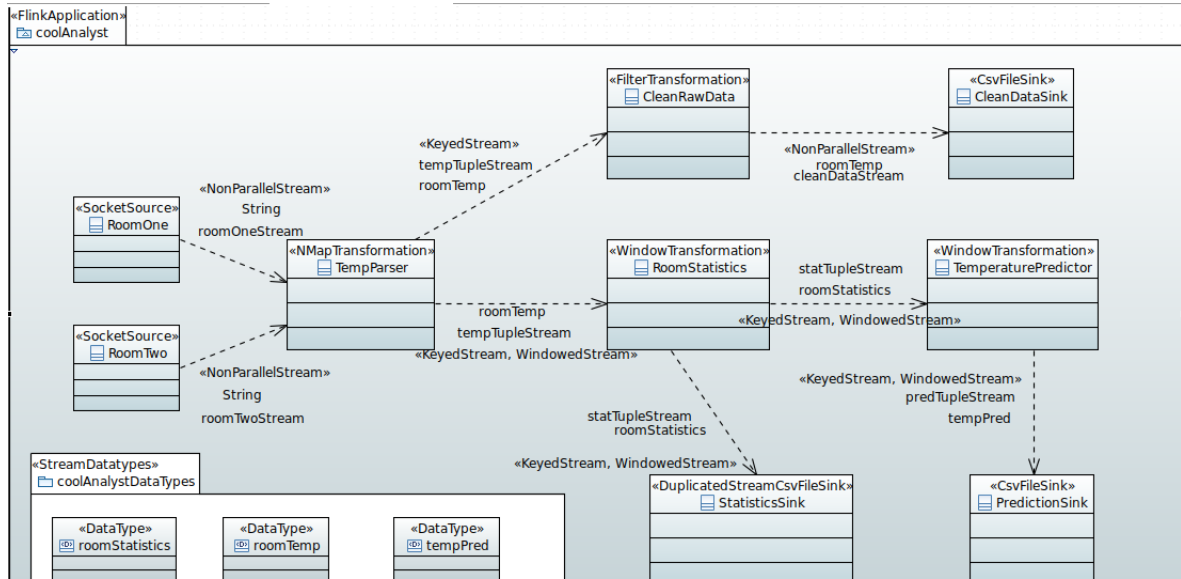


Figure 5.8: Cool Analyst StreamGen DIA

the characteristics of the protected stream. This is why only one stream requires to be protected by privacy policies.

In order to do this, the PrivacyProtectingStream stereotype is applied in the stream known as tempTupleStream flowing from the TempParser transformation to the RoomStatistics transformation. Once the stereotype is applied, the properties for such stereotype are defined. In this case, the stream is protected by a DSEP (protectedByDSEP is true) but it is not protected by any VCP (protectedByVCP is false). Moreover, protectedStreamConf data type is defined with the same values that the protected streams in the Great Seller example were designed. In the table 5.6 can be seen how are defined all these values.

Property	Field Name	Field Value
protectedByVCP	-	false
protectedByDSEP	-	true
protectedStreamConf	monitoringActive	false
	timestampServerIp	
	timeStampServerPort	-1
	topologyParallelism	1
	simulateRealisticScenario	false
	allowedLateness	0
	logDir	/home/cablan/Desktop/thesis/conf/

Table 5.6: PrivacyProtectingStream Cool Analyst TempTupleStream

The value for the timestampServerIp is empty because, as in the previous example, it is defined to nothing. All these values can be modified but they are considered as the default values for any protected stream.

In the figure 5.9 can be seen the applied stereotype in the non-privacy-aware application in order to defined which streams are protected.




Figure 5.9: Cool Analyst Privacy Aware StreamGen DIA

5.2.4.2 Privacy External Sources Definition

After defining the protected streams, the external privacy sources are introduced. In this case, the predefined DSEP must be applied every time that an employee who is not the manager of the Panettone 4.0 factory tries to see the statistics of the first chamber. Then, a TXT file with all the employee identifiers and assuming that they try to access with analytical purposes is created and accessed by the SCV source.

In order to implement this, first of all, a package node called ExternalPrivacySources is introduced in the application model and the PrivacyPolicyPackage stereotype is applied on it in order to make more intuitive how to introduce the SCV source and the privacy rules source.

After that, a class node called SCVSource is introduced in the ExternalPrivacySources package and the PrivContTextFileSource stereotype is applied on it. This stereotype requires to define the pathToFile (string) property, in the table 5.7 can be seen such path. Moreover, in the table 5.8 can be seen the content of the text file.

Once the SCV source is completely defined, the privacy rules source is introduced. For such purpose, a class node called PrivacyRulesSource is introduced in the ExternalPrivacySources package. In this class the PrivPolYamlFileSource is applied and its property, pathToFile (string), is defined according to the table 5.7.

Class	Property Name	Property Value
SCVSource	pathToFile	/home/cablan/Desktop/thesis/conf/
PrivacyRulesSource	pathToFile	/home/cablan/Desktop/thesis/conf/

Table 5.7: Cool Analyst External Sources Properties

At this point, the privacy-aware DIA is implemented and the configuration is run in order to generate the Java codes of the application. In the appendix C can be seen how to run such configuration.

Observer Identifier	Purpose	Role
employee1	analytics	employee
employee2	analytics	employee
employee3	analytics	employee

Table 5.8: Cool Analyst SCV Text File

5.3 Discussion

Great Seller and Cool Analyst are two DIA developed in this document in order to evaluate the privacy approach explained in chapter 4. The aim of this evaluation is to answer some previously stated questions:

- Can software developers with few knowledge about privacy policies apply this approach?

As it can be seen in both examples, software developers just need to know if a stream is protected by VCPs and DSEPs independently of the number of policies that are applied on them. Moreover, software developers apply the privacy policies by adding a stereotype to the stream that must be protected in the same way that a windowed or keyed stream is specified.

- Can software developers with few knowledge about Flink language apply this approach?

Due to the fact that the approach imports an already developed library with a predefined structure, the generated Flink codes by means of Acceleo follow the same structure independently of the implemented DIA. This is why, codes must not be modified by software developers and, then, it is not required that they need any knowledge about the target language, in this case Flink language.

- Is easily applicable the privacy approach once the non-privacy-aware DIA is already implemented?

In both case studies developed in this chapter, the privacy-aware DIA is reached only by following two simple steps (protected streams definition and privacy external sources definition) once the non-privacy-aware application is already implemented. Moreover, both steps are defined considering the already existing approach to introduce data types by means of a package and by applying a stereotype in those streams that must be protected.

In summary, with this two case studies can be seen how the objectives stated in chapter 3 are satisfied.

Chapter 6

Conclusions

This document faces three goals in order to develop a language which is able to apply privacy policy in streaming applications after compiling the language by means of an Acceleo compiler. Such goals are:

- To extend the StreamGen high level modeling approach to enable the modeling of privacy-aware streaming applications.
- To allow software developers with few knowledge about privacy policies to handle the design of privacy-aware streaming applications.
- From a non-privacy-aware application reach a privacy-aware application easily.

This goals are satisfied after evaluating two case studies but in spite of the privacy approach developed along this document allows to easily apply privacy policy in streaming applications by means of a metamodeling design, it takes into account that privacy policies are introduced in the DIA by means of a predefined YAML. This predefined YAML must be developed by someone with some knowledge about privacy policies and who is capable to understand the privacy requirements of the application users. This is why, it is stated as a future work to develop a metamodel which is able to generate the YAML file. Such metamodel should be filled by users in order to reduce the misunderstood that can appear when the user specifies his/her privacy requirements to a intermediary, for example a software developer.

Bibliography

- [1] Computing History Web Page
<https://www.livescience.com/20718-computer-history.html>
- [2] A.M. Middleton. *Data-Intensive Technologies for Cloud Computing*. Handbook of Cloud Computing. Springer, 2010.
- [3] Interactive and Batch Processing
<https://www2.le.ac.uk/offices/itservices/ithelp/services/hpc/run-computation/interactive-or-batch>
- [4] Ververica Web Page
<https://www.ververica.com/what-is-stream-processing>
- [5] General Data Protection Regulation Web Page
<https://gdpr-info.eu/>
- [6] Michele Guerriero, Damian Andrew Tamburri, Elisabetta Di Nitto. *Defining, Enforcing and Checking Privacy Policies In Data-Intensive Applications*. 13th International Symposium on Software Engineering for Adaptive and Self-Managing Systems, 2018.
- [7] D.B. Skillicorn and D. Talia. *Models and languages for parallel computation*. ACM Computing Surveys, Vol. 30, No. 2, 1998, pp. 123-169.
- [8] I. Gorton, P. Greenfield, A. Szalay and R. Williams. *Computing in the 21st Century*. IEEE Computer, Vol. 41, No. 4, 2008, pp. 30-32.
- [9] W.E. Johnston. *High-Speed, Wide Area, Data Intensive Computing: A Ten Year Retrospective*. IEEE Computer Society, 1998.
- [10] M. Gokhale, J. Cohen, A. Yoo and W.M. Miller. *IEEE: Hardware Technologies for High-Performance Data-Intensive Computing*. IEEE Computer, Vol. 41, No. 4, 2008, pp. 60-68.
- [11] D. Ravichandran, P. Pantel and E. Hovy. *The terascale challenge*. Proceedings of the KDD Workshop on Mining for and from the Semantic Web, 2004.
- [12] J. Gray. *Distributed Computing Economics*. ACM Queue, Vol. 6, No. 3, 2008, pp. 63-68.
- [13] R.E. Bryant. *Data Intensive Scalable Computing*. 2008.
- [14] Hadoop Tutorial
<https://data-flair.training/blogs/hadoop-tutorial/>

- [15] Spark Tutorial
<https://data-flair.training/blogs/spark-tutorial/>
- [16] Apache Flink Tutorial
<https://data-flair.training/blogs/apache-flink-tutorial/>
- [17] Data Privacy Blog
<https://www.varonis.com/blog/data-privacy/>
- [18] Privacy Policy Definition
<https://www.techopedia.com/definition/23679/privacy-policy>
- [19] Ken Thompson. *Reflections on Trusting Trust*. Communications of the ACM, 1984.
- [20] Health Insurance Portability and Accountability Act Blog
<https://www.varonis.com/blog/hipaa-compliance/>
- [21] Gramm-Leach-Bliley Act Blog
<https://www.varonis.com/blog/we-need-to-talk-about-gramm-leach-bliley-glb-the->
- [22] California Consumer Privacy Act Blog
<https://www.varonis.com/blog/california-consumer-privacy-act-ccpa/>
- [23] Caroline Dewi Puspa Kencana Ramli, Hanne Riis Nielson, Flemming Nielson. *The logic of XACML*.
- [24] Markus Lorch, Seth Proctor, Rebekah Lepro, Dennis Kafura, Sumit Shah. *First Experiences Using XACML for Access Control in Distributed Systems*.
- [25] Modeling Tutorial
<https://www.transentis.com/methods-techniques/models-and-metamodels/>
- [26] Dimitris Karagiannis, Harald Kühn *Metamodeling Platforms* Springer-Verlag, 2002
- [27] Unified Modeling Language Tutorial
<https://www.visual-paradigm.com/guide/uml-unified-modeling-language/what-is-uml/>
- [28] Temperature Control For Yeast Fermentation
http://www.theartisan.net/temperature_control_baking_1.htm

Appendices

Appendix A

Papyrus Project Creation

In order to implement the application, first of all, a new project has to be created in Eclipse. This project is going to be a Papyrus project (figure A.1).

The first field that has to be specified is the Architecture Context. Here, the default values are the right ones as an UML model is what is required to implement with Papyrus (figure A.2).

After that, the name of the project has to be defined. Due to the fact that it is going to implement an application called Great Seller, the name of the Papyrus project is going to be greatSellerApp (figure A.3).

The next step is to select the representation kind and to choose the StreamGen profile. StreamGen requires Papyrus to make a class diagram representation and, as the StreamGen project is already in the Eclipse workspace, the profile can be found browsing in the workspace (streamngen/profile/StreamUML.profile.uml). As it can be seen in the figure A.4.

Finally, we can finish with the new Papyrus project initialization (figure A.5).




Figure A.1: New Papyrus Project




Figure A.2: Architecture Context

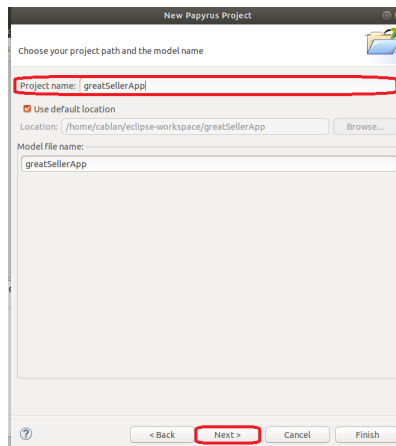


Figure A.3: Project Name Definition




Figure A.4: Representation Kind

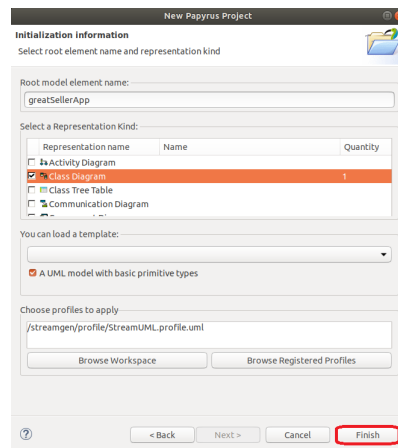


Figure A.5: Papyrus Project Definition

Appendix B

Support Python Codes

B.1 Great Seller Stock Creation

```
# Libraries
import random, shelve

# Predefined variables
numberProducts = 25 # Great Seller number of products available
priceLowerLimit = 1
priceUpperLimit = 200
greatSellerStock = []

# Binary shelf file creation
shelfFile = shelve.open('stocks')

# Creating the stock
for eachProduct in range(numberProducts):
    name = 'product' + str(eachProduct + 1)
    price = random.randint(priceLowerLimit, priceUpperLimit)
    product = {'name': name, 'price': price}
    greatSellerStock = greatSellerStock + [product]

# Writing the stock in the binary shelf file
shelfFile['greatSellerStock'] = greatSellerStock

# Closing the binary shelf file
shelfFile.close()

# Printing the generated stock in Latex table format
string = ''

for eachProduct in greatSellerStock:
    string = string + eachProduct['name'] + '&' + str(eachProduct['price'])
    string = string + '\hline' + '\n'

print(string)
```

B.2 Great Seller Input Tuples Creation

```
# Libraries
import random, sys, shelve

# Predefined variables
dataSubjectList = [ 'Bob', 'Carlos', 'Elisabetta', 'Michele' ]

if len(sys.argv) < 2:
    print( 'Usage: _python3 _tupleCreator.py _[fileLength] ' )
    sys.exit()

# Running defined variables
fileLength = int(sys.argv[1])

# Open the source binary shelf file
shelfFile = shelve.open('stocks')
greatSellerStock = shelfFile['greatSellerStock']

# Open the destination file in write mode
fileName = 'tuples' + str(fileLength) + 'File.txt'
tuplesFile = open(fileName, 'w')

# In each iteration we create a random tuple
for eachValue in range(fileLength):

    # Tuple values creation
    transactionId = str(eachValue + 1)
    dataSubject = dataSubjectList[random.randint(0, len(dataSubjectList))]
    product = greatSellerStock[random.randint(0, len(greatSellerStock))]
    amount = str(product['price'])
    recipientId = product['name']

    # Tuple List creation
    tupleValues = [transactionId, dataSubject, amount, recipientId]

    # Tuple joined in a string separated by a comma
    builtTuple = ','.join(tupleValues) + '\n'

    # Tuple written in the file
    tuplesFile.write(builtTuple)

# Close the destination file
tuplesFile.close()
```

B.3 Great Seller Input Server

```
import socket, time
```

```
# Predefined variables
waitingTime = 1
fileLength = 1000

# Open the source file in read mode
fileName = 'tuples' + str(fileLength) + 'File.txt'
fileName = '/home/cablan/Desktop/thesisFiles/' + fileName
tuplesFile = open(fileName, 'r')

# Create a socket
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)

# Define the host and the port and listen into them
host = 'localhost'
port = 9999
s.bind((host, port))
s.listen(5)

while True:
    # Accept the requested connection
    clientSocket, addr = s.accept()
    print("Connection_accepted_from_" + repr(addr[1]))

    # Send each tuple contained in the file
    for eachTuple in tuplesFile:
        time.sleep(waitingTime)
        print('Sending_tuple:_ ' + '(' + eachTuple[0:-1] + ')')
        clientSocket.send(bytes(eachTuple, 'utf-8'))

    # Close the connection
    clientSocket.close()
```

B.4 Cool Analyst Temperature Chamber 1 Model

```
import random, math

import matplotlib.pyplot as plt

def sin(upperBoundary, lowerBoundary, freq, initialDelta):
    # Fix Variables
    midValue = lowerBoundary + (upperBoundary - lowerBoundary) / 2
    angularSpeed = 2 * math.pi * freq
    y = []

    for time in range(int(1/freq)):
        temp = midValue + ((upperBoundary - lowerBoundary) / 2) * math.sin
```

```

        y += [temp]

    return y

roomId = 'room1'
timeLength = 100

# Creating Plotting Lists
x = []
y = []

# Temperature Limits and Average
lowerLimit = 22
upperLimit = 29
avgTemp = lowerLimit + (upperLimit - lowerLimit) / 2

# Initial Time
time = 0

# Generating the function
while True:

    for staticValues in range(10):
        y += [avgTemp]

    randValue = random.randint(0,1)

    if randValue == 0:
        ysin = sin(avgTemp, lowerLimit, 1/16, math.pi / 2)
    else:
        ysin = sin(upperLimit, avgTemp, 1/28, math.pi * (-1) / 2)

    y += ysin
    time += 1
    staticValues = 0

    if len(y) > timeLength:
        break

for i in range(len(y)):
    x += [i]

fileObject = open(' /home/cablan/Desktop/thesisFiles/inputs/room1Sin.txt'
for eachTemp in y:
    fileObject.write(roomId + ', ' + str(eachTemp) + '\n')

# Creating the plot

```

```
# plotting the points
plt.plot(x, y)

# naming the x axis
plt.xlabel('Time_(s)')
# naming the y axis
plt.ylabel('Temperature_(Celsius)')

# giving a title to my graph
plt.title('Chamber_1')

# function to show the plot
plt.show()
```

B.5 Cool Analyst Temperature Chamber 2 Model

```
import random, math

import matplotlib.pyplot as plt

def sin(upperBoundary, lowerBoundary, freq, initialDelta):
    # Variables
    midValue = lowerBoundary + (upperBoundary - lowerBoundary) / 2
    angularSpeed = 2 * math.pi * freq
    y = []

    for time in range(int(1/freq)):
        temp = midValue + ((upperBoundary - lowerBoundary) / 2) * math.sin(
            initialDelta + angularSpeed * time)
        y += [temp]

    return y

roomId = 'room2'
timeLength = 120

# Creating Plotting Lists
x = []
y = []

# Temperature Limits and Average
lowerLimit = 20
upperLimit = 24
avgTemp = lowerLimit + (upperLimit - lowerLimit) / 2

# Initial Time
time = 0
```

```

# Generating the function
while True:

    for staticValues in range(30):
        y += [avgTemp]

    randValue = random.randint(0,1)

    if randValue == 0:
        ysin = sin(avgTemp, lowerLimit, 1/16, math.pi / 2)
    else:
        ysin = sin(upperLimit, avgTemp, 1/28, math.pi * (-1) / 2)

    y += ysin
    time += 1
    staticValues = 0

    if len(y) > timeLength:
        break

for i in range(len(y)):
    x += [i]

fileObject = open('/home/cablan/Desktop/thesisFiles/inputs/room2Sin.txt', 'a')
for eachTemp in y:
    fileObject.write(roomId + ', ' + str(eachTemp) + '\n')

# Creating the plot
# plotting the points
plt.plot(x, y)

# naming the x axis
plt.xlabel('Time_(s)')
# naming the y axis
plt.ylabel('Temperature_(Celsius)')

# giving a title to my graph
plt.title('Chamber_2')

# function to show the plot
plt.show()

```

B.6 Cool Analyst Chamber 1 Server

```
import socket, time, sys
```



```
# Predefined variables
waitingTime = 1

# Choose fileLength
if len(sys.argv) == 2:
    fileLength = sys.argv[1]
else:
    fileLength = 10

# Open the source file in read mode
fileName = 'room1.txt'
fileName = '/home/cablan/Desktop/thesisFiles/inputs/' + fileName
tuplesFile = open(fileName, 'r')

# Create a socket
s = socket.socket(socket.AF_INET, socket.SOCKSTREAM)

# Define the host and the port and listen into them
host = 'localhost'
port = 8888
s.bind((host, port))
s.listen(5)

# Accept the requested connection
clientSocket, addr = s.accept()
print("Connection_accepted_from_" + repr(addr[1]))

time.sleep(5)

# Send each tuple contained in the file
for eachTuple in tuplesFile:
    print('Sending_tuple:_ ' + '(' + eachTuple[0:-1] + ')')
    clientSocket.send(bytes(eachTuple, 'utf-8'))
    time.sleep(waitingTime)

# Close the connection
clientSocket.close()

# Close the file
tuplesFile.close()
```

B.7 Cool Analyst Chamber 2 Server

```
import socket, time, sys

# Predefined variables
waitingTime = 1
```

```
# Choose fileLength
if len(sys.argv) == 2:
    fileLength = sys.argv[1]
else:
    fileLength = 10

# Open the source file in read mode
fileName = 'room2.txt'
fileName = '/home/cablan/Desktop/thesisFiles/inputs/' + fileName
tuplesFile = open(fileName, 'r')

# Create a socket
s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)

# Define the host and the port and listen into them
host = 'localhost'
port = 8887
s.bind((host, port))
s.listen(5)

# Accept the requested connection
clientSocket, addr = s.accept()
print("Connection_accepted_from_" + repr(addr[1]))

time.sleep(5)

# Send each tuple contained in the file
for eachTuple in tuplesFile:
    print('Sending_tuple:_ ' + '(' + eachTuple[0:-1] + ')')
    clientSocket.send(bytes(eachTuple, 'utf-8'))
    time.sleep(waitingTime)

# Close the connection
clientSocket.close()

# Close the file
tuplesFile.close()
```

Appendix C

Run Configuration

In order to generate the Java codes from the Papyrus model, first of all we need to create a new Java project in Eclipse (File -> New -> Java Project). This project is going to be called 'greatSellerCodes' as it can be seen in the figure C.1.

Once the Java project is created, its structure has to be modified in order to allow to be compatible with Maven. As the figure C.2 shows, firstly the Java project is composed of a JRE System Library and a src folder. This src folder has to be deleted, leaving the Java project only with the JRE System Library as it can be seen in the figure C.3.

The next step is to create a source folder with a predefined structure as it is going to be specified now. Firstly, in the properties for the project, in the Java Build Path field, a folder has to added as the figure C.4 specifies. In the default output folder a specific structure has to be written, greatSellerCodes/src/main/java. After this is written, the changes have to be applied and then click on 'Add Folder...'. .

In order to create the source folder, the java folder has to be selected and then click on 'OK' as it is shown in figure C.5. Then the changes have to be applied and, then, the window Properties for greatSellerCodes has to be closed (figure C.6).

After these steps, the project folder should look like it is shown in the figure C.7 at the Package Explorer. And then, the project is ready to convert it into a Maven project. In order to do this, going to the 'Configure' option of the project and then clicking on 'Convert to Maven Project'. Then, a POM file has to be created and it




Figure C.1: Java Project Creation




Figure C.2: Java Project Initial Structure




Figure C.3: Java Project Final Structure




Figure C.4: Maven Project Structure Creation




Figure C.5: Maven Project Source Folder Creation

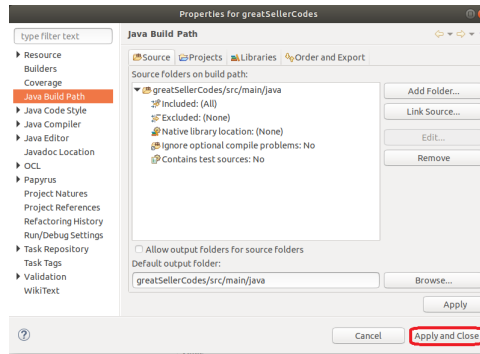


Figure C.6: Maven Project Properties

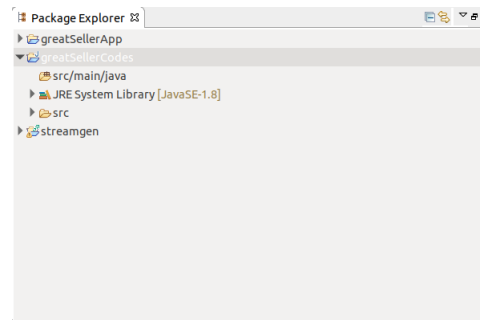


Figure C.7: Maven Project Structure

is going to be created the default one as it can be seen in the figure C.8. At this point, the structure of the project has to be the same that the one shown in the figure C.9.

Now it is time to run the configuration (Run -> Run Configurations...). In the figure C.10 is shown how this window has to be filled.




Figure C.8: POM File Creation

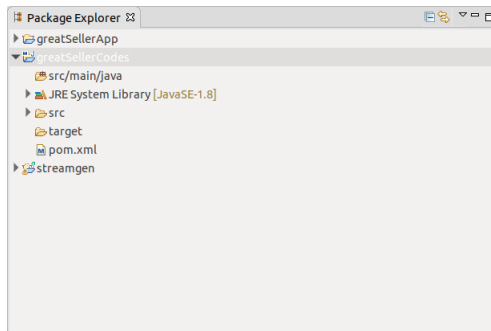


Figure C.9: Final Maven Project

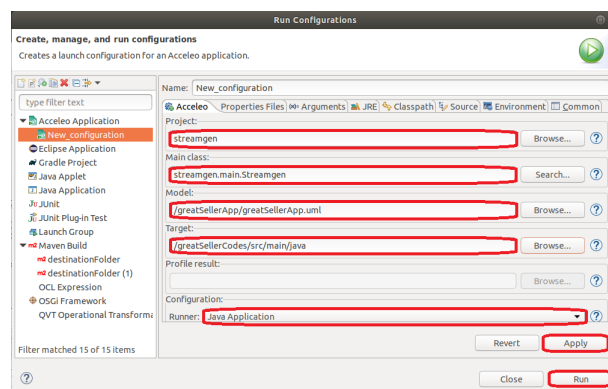


Figure C.10: Run Configuration